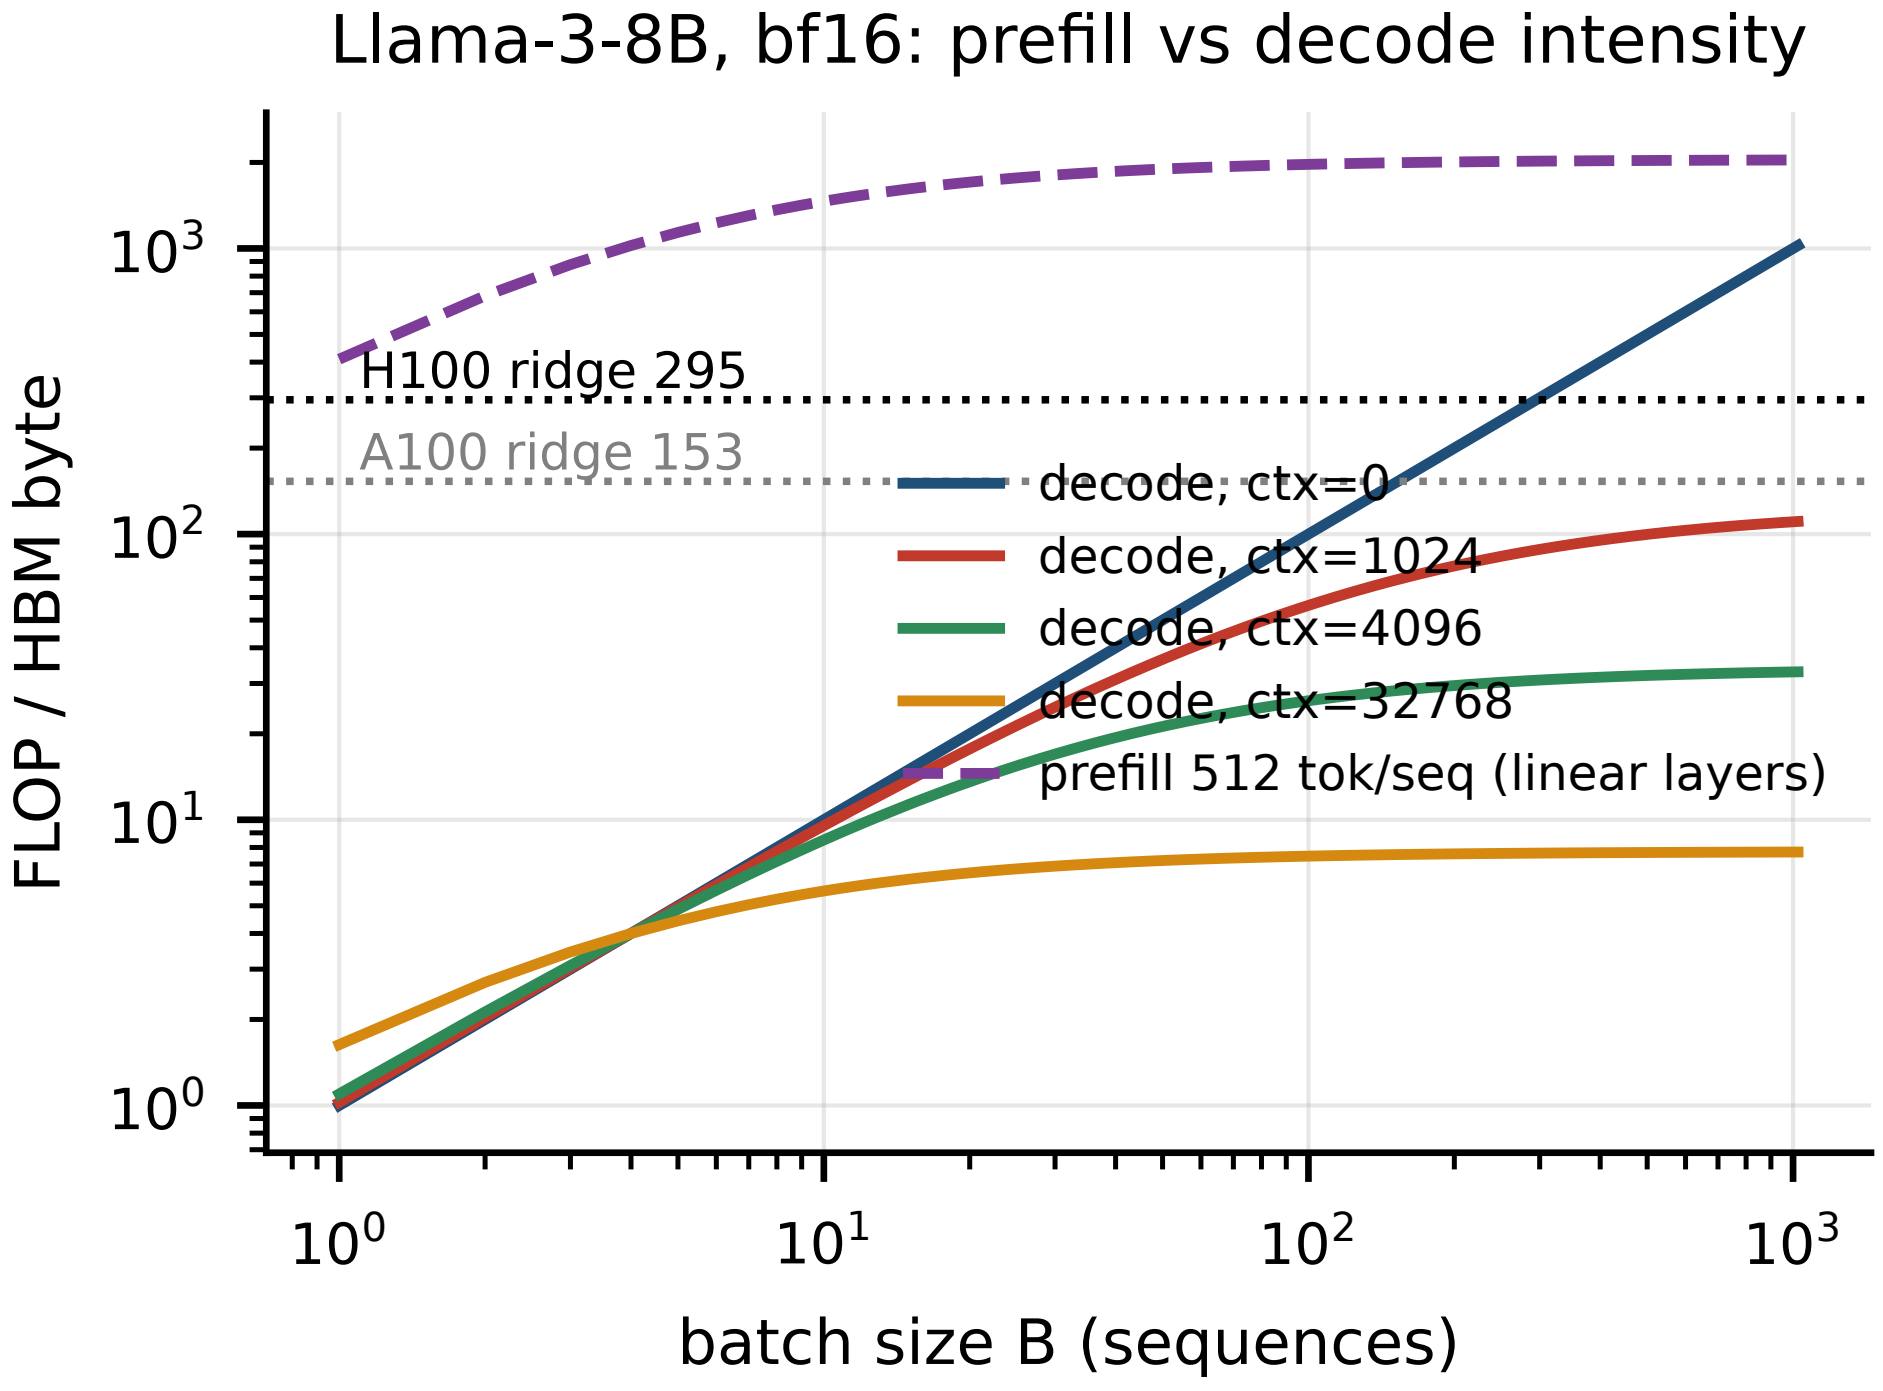
Llama-3-8B, bf16: prefill vs decode intensity
FLOP / HBM byte
batch size B (sequences)
H100 ridge 295
A100 ridge 153
decode, ctx=0
decode, ctx=1024
decode, ctx=4096
decode, ctx=32768
prefill 512 tok/seq (linear layers)
$10^0$
$10^1$
$10^2$
$10^3$