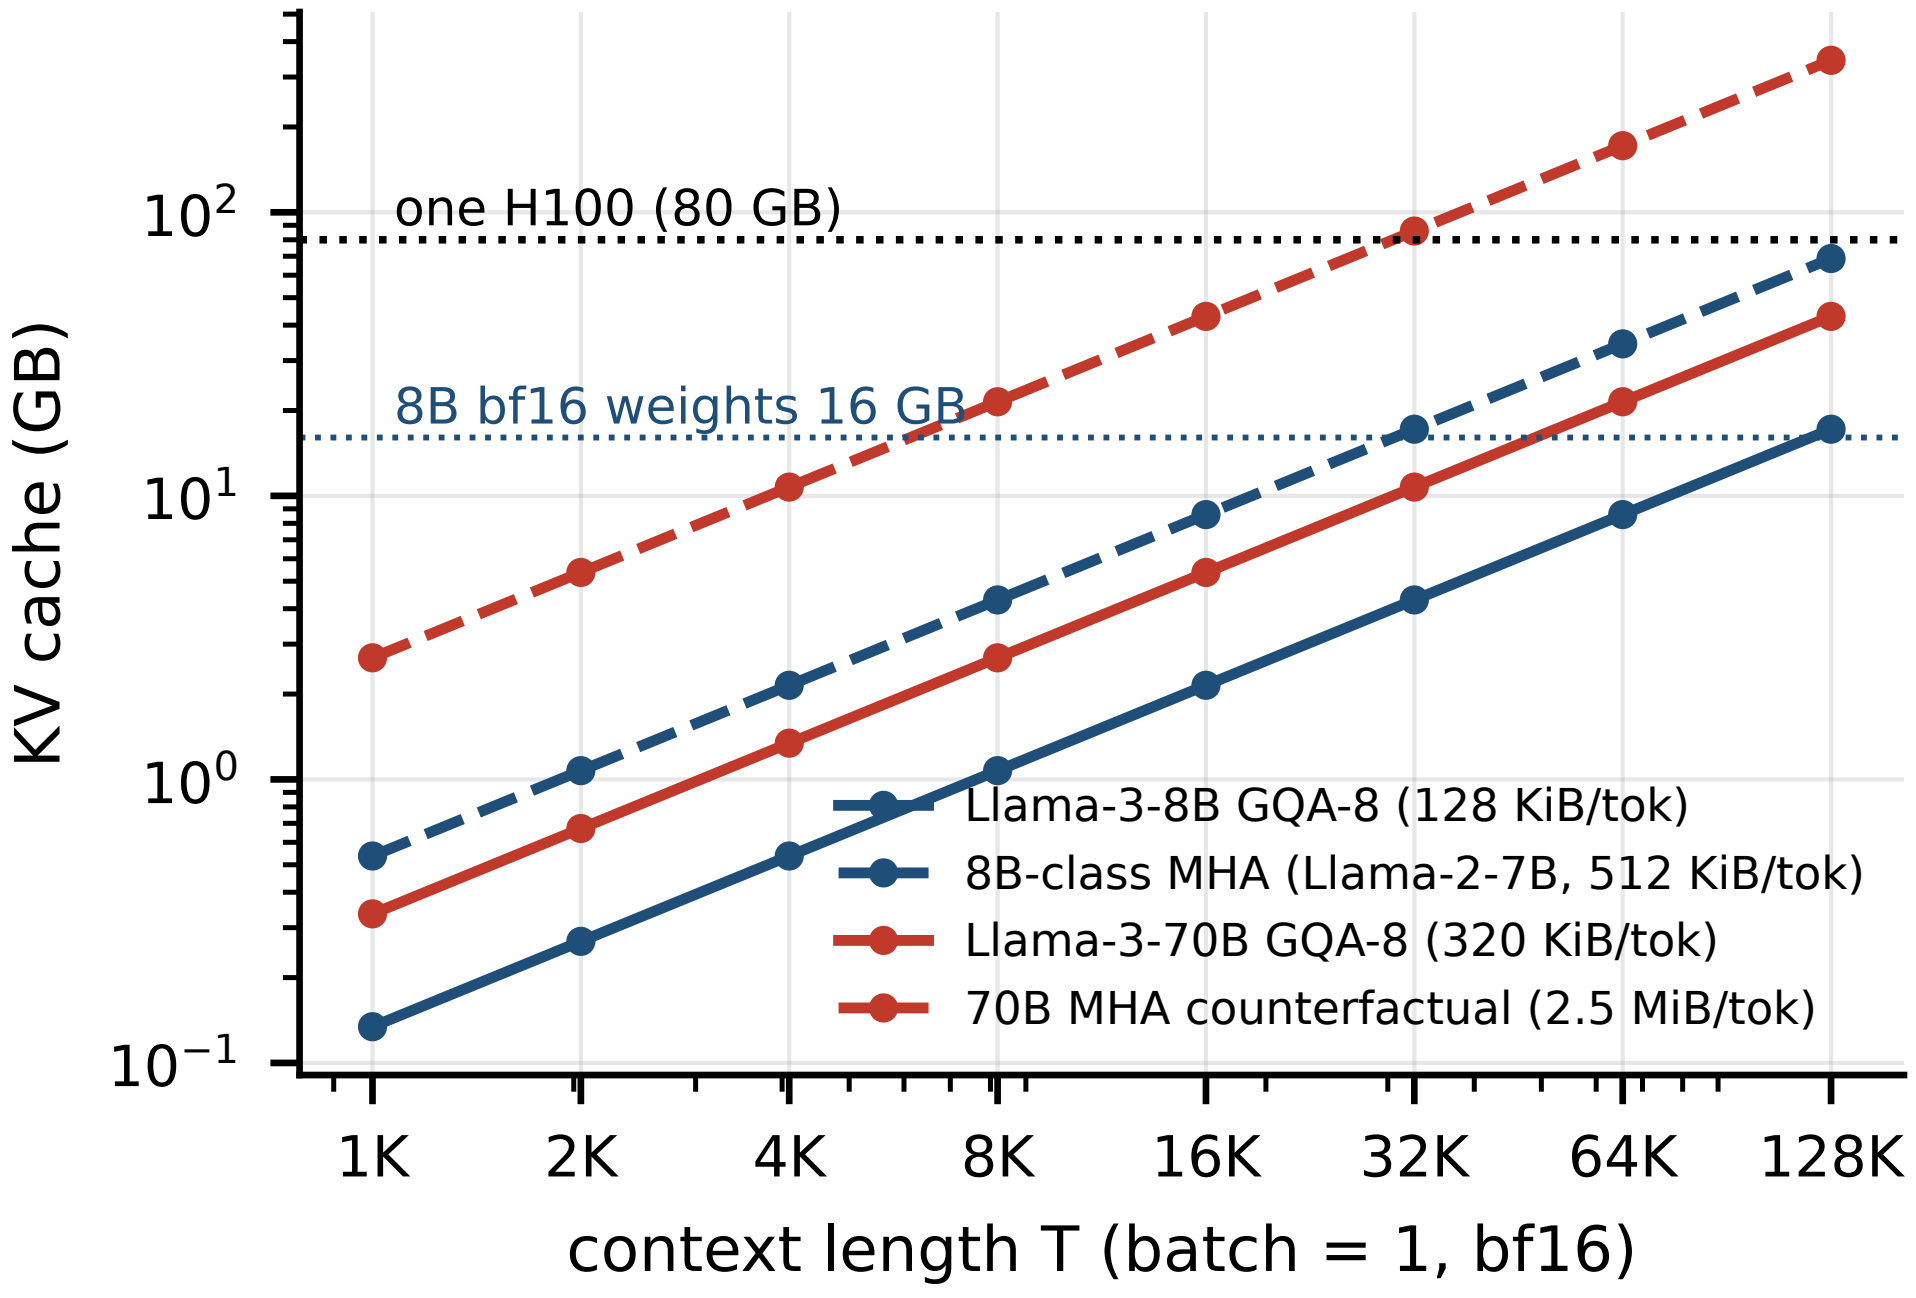
one H100 (80 GB)
8B bf16 weights 16 GB
Llama-3-8B GQA-8 (128 KiB/tok)
8B-class MHA (Llama-2-7B, 512 KiB/tok)
Llama-3-70B GQA-8 (320 KiB/tok)
70B MHA counterfactual (2.5 MiB/tok)
KV cache (GB)
10^2
10^1
10^0
10^{-1}
1K
2K
4K
8K
16K
32K
64K
128K
context length T (batch = 1, bf16)