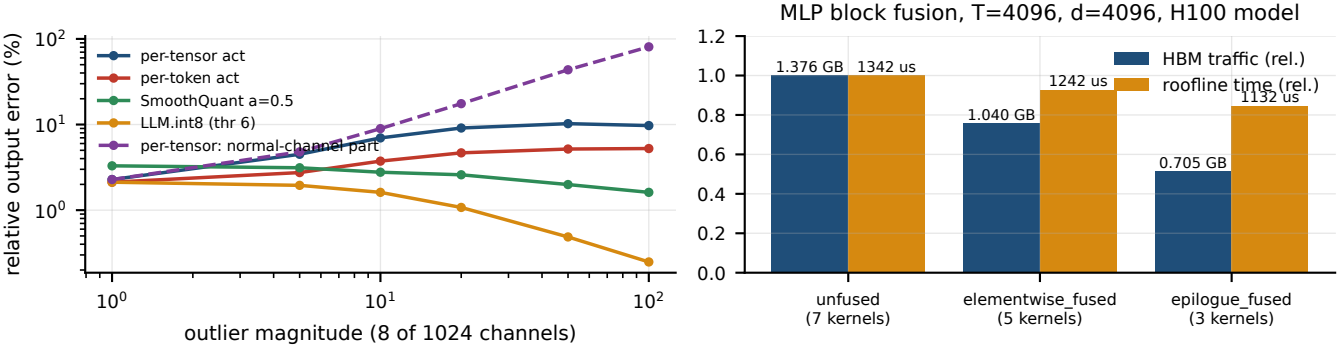
relative output error (%)
10^2
10^1
10^0
per-tensor act
per-token act
SmoothQuant a=0.5
LLM.int8 (thr 6)
per-tensor: normal-channel part
10^0
10^1
10^2
outlier magnitude (8 of 1024 channels)
MLP block fusion, T=4096, d=4096, H100 model
1.2
1.0
0.8
0.6
0.4
0.2
0.0
HBM traffic (rel.)
roofline time (rel.)
1.376 GB
1342 us
1.040 GB
1242 us
0.705 GB
1132 us
unfused
(7 kernels)
elementwise_fused
(5 kernels)
epilogue_fused
(3 kernels)