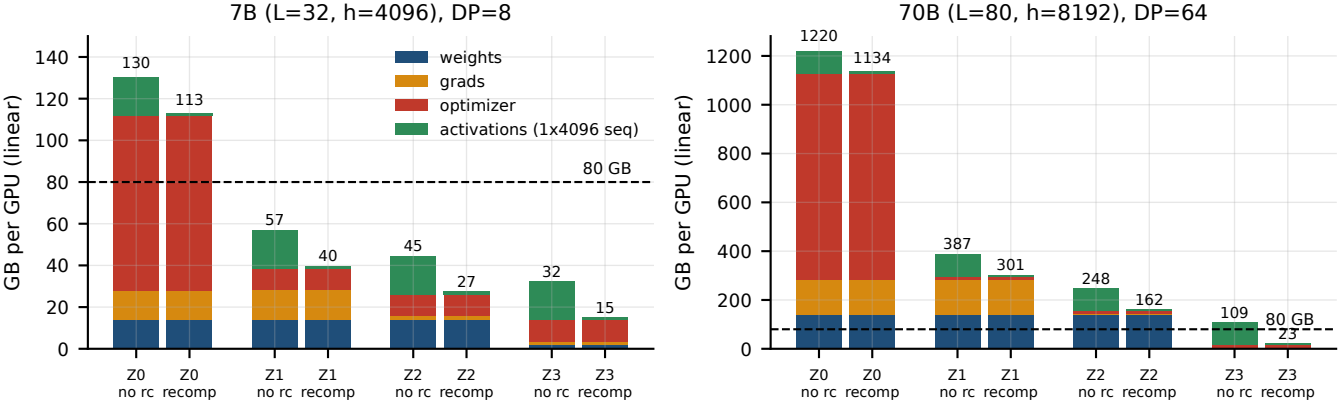
7B (L=32, h=4096), DP=8
GB per GPU (linear)
weights
grads
optimizer
activations (1x4096 seq)
80 GB
130
113
57
40
45
27
32
15
Z0 no rc
Z0 recomp
Z1 no rc
Z1 recomp
Z2 no rc
Z2 recomp
Z3 no rc
Z3 recomp
70B (L=80, h=8192), DP=64
1220
1134
387
301
248
162
109
23