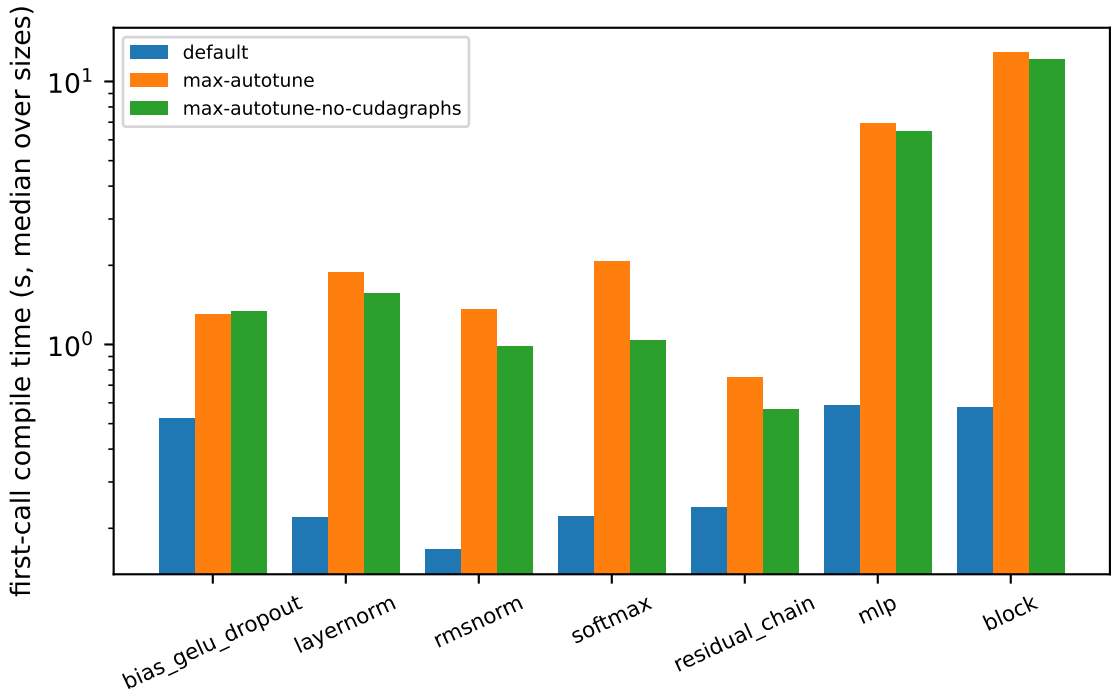
default
max-autotune
max-autotune-no-cudagraphs
first-call compile time (s, median over sizes)
10^1
10^0
bias_gelu_dropout
layernorm
rmsnorm
softmax
residual_chain
mlp
block

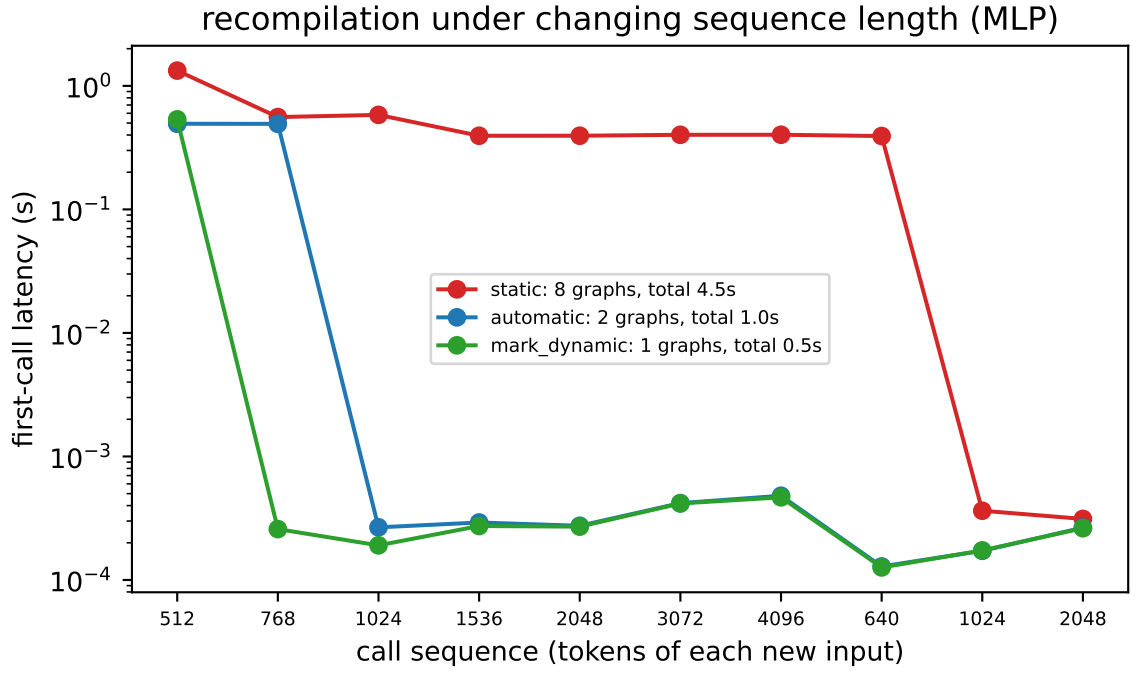
recompilation under changing sequence length (MLP)
static: 8 graphs, total 4.5s
automatic: 2 graphs, total 1.0s
mark_dynamic: 1 graphs, total 0.5s
first-call latency (s)
10^0
10^{-1}
10^{-2}
10^{-3}
10^{-4}
512
768
1024
1536
2048
3072
4096
640
1024
2048
call sequence (tokens of each new input)