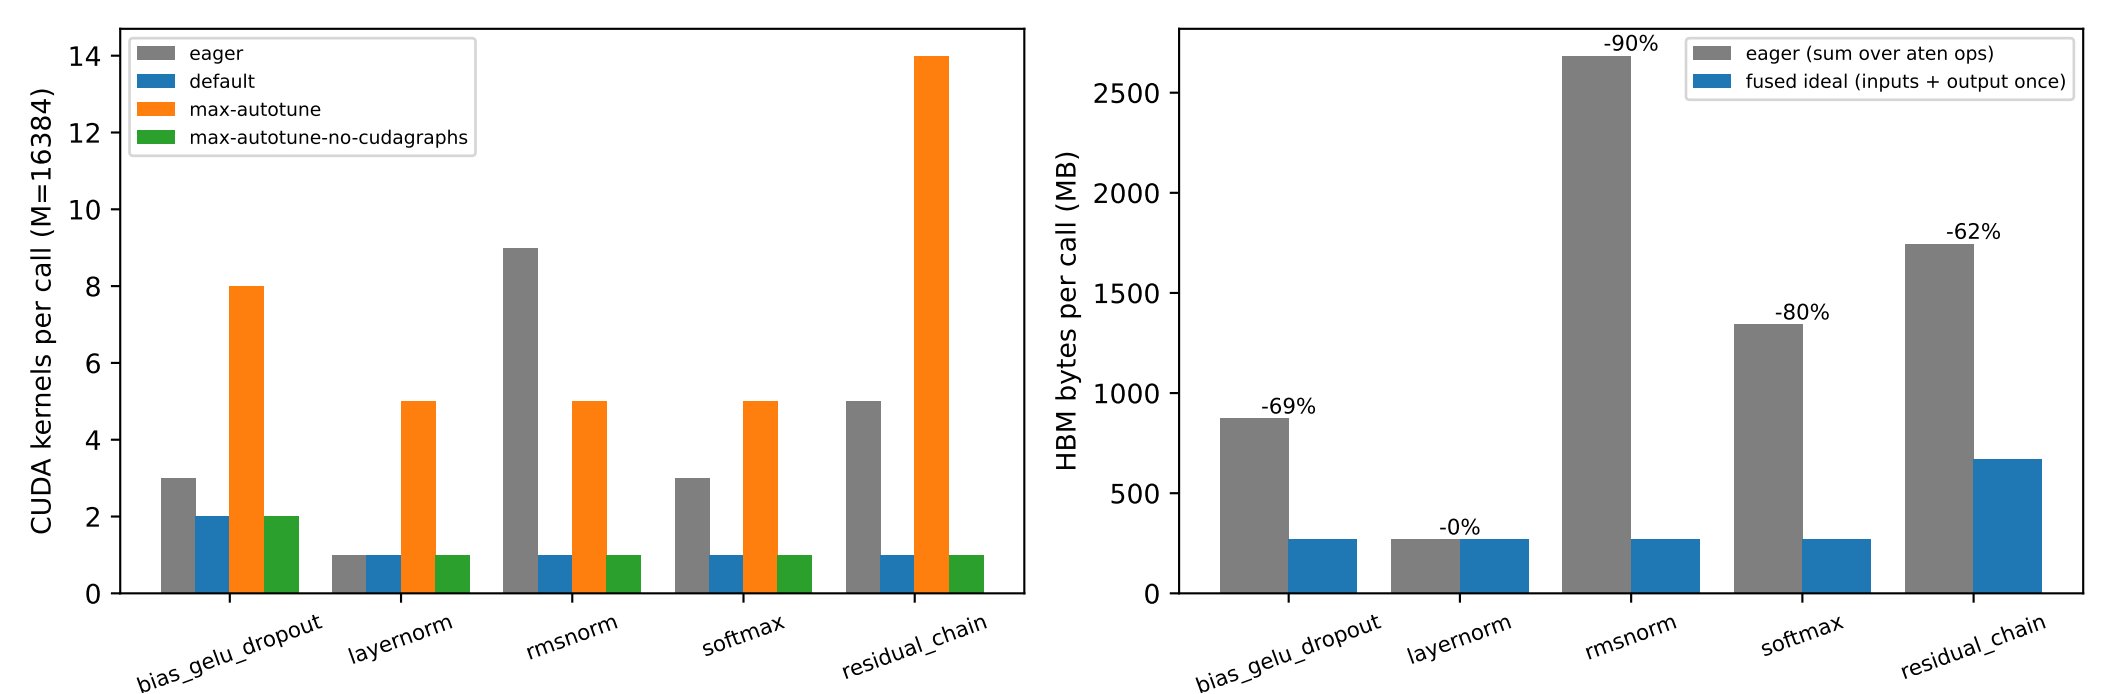
CUDA kernels per call (M=16384)
eager
default
max-autotune
max-autotune-no-cudagraphs
0
2
4
6
8
10
12
14
bias_gelu_dropout
layernorm
rmsnorm
softmax
residual_chain
HBM bytes per call (MB)
eager (sum over aten ops)
fused ideal (inputs + output once)
0
500
1000
1500
2000
2500
-69%
-0%
-90%
-80%
-62%
bias_gelu_dropout
layernorm
rmsnorm
softmax
residual_chain