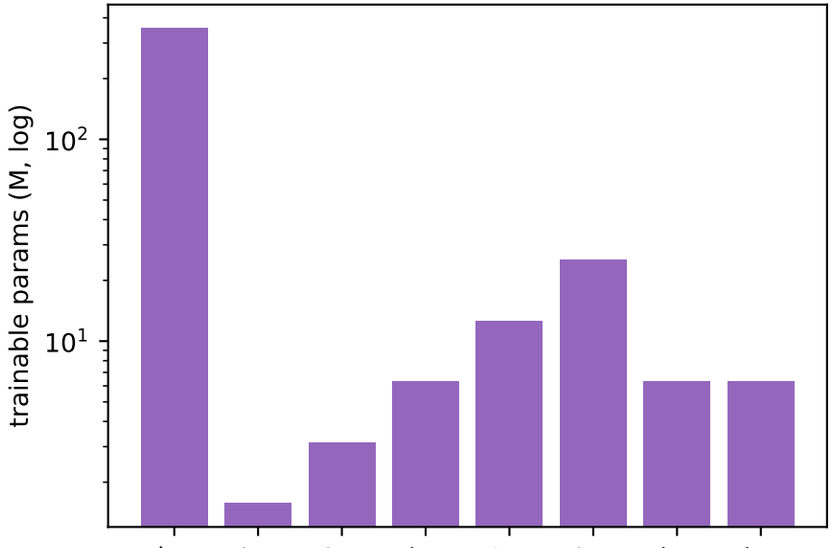
trainable params (M, log)
10^2
10^1
full
lora_r4
lora_r8
lora_r16
lora_r32
lora_r64
qlora_int8_r16
qlora_int4_r16

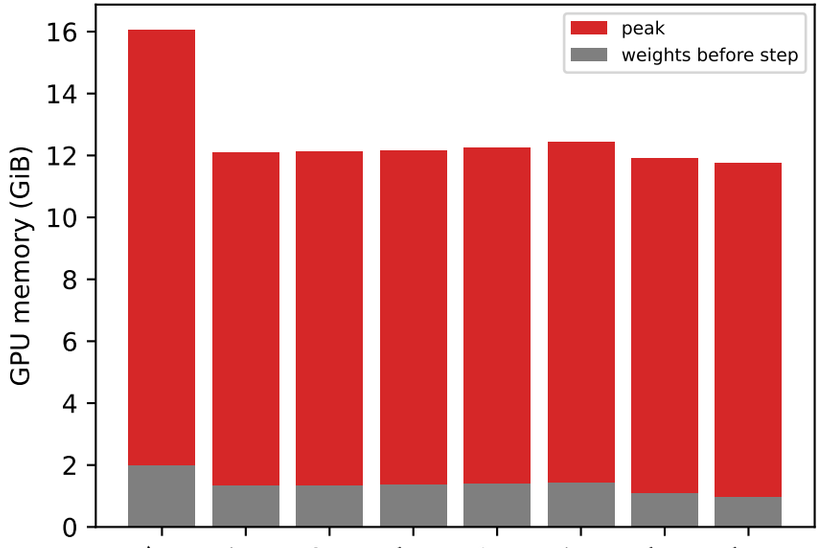
GPU memory (GiB)
peak
weights before step
0
2
4
6
8
10
12
14
16
full
lora_r4
lora_r8
lora_r16
lora_r32
lora_r64
qlora_int8_r16
qlora_int4_r16

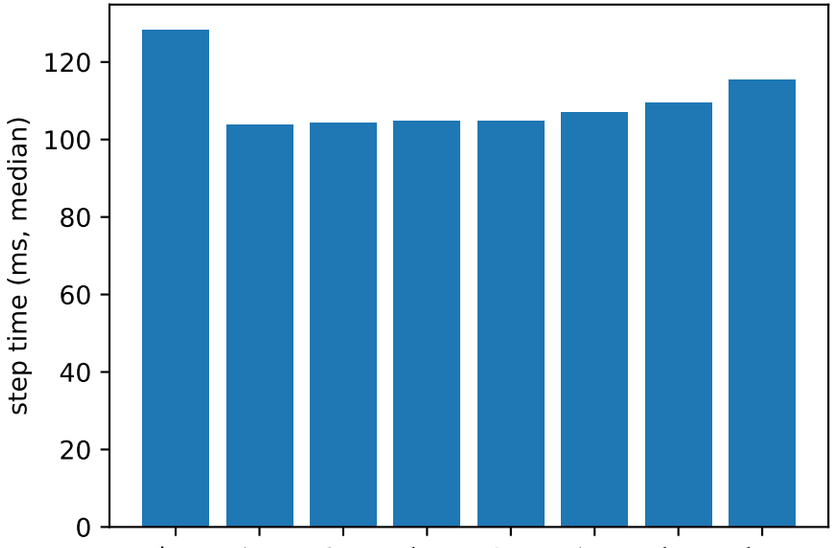
step time (ms, median)
0
20
40
60
80
100
120
full
lora_r4
lora_r8
lora_r16
lora_r32
lora_r64
qlora_int8_r16
qlora_int4_r16