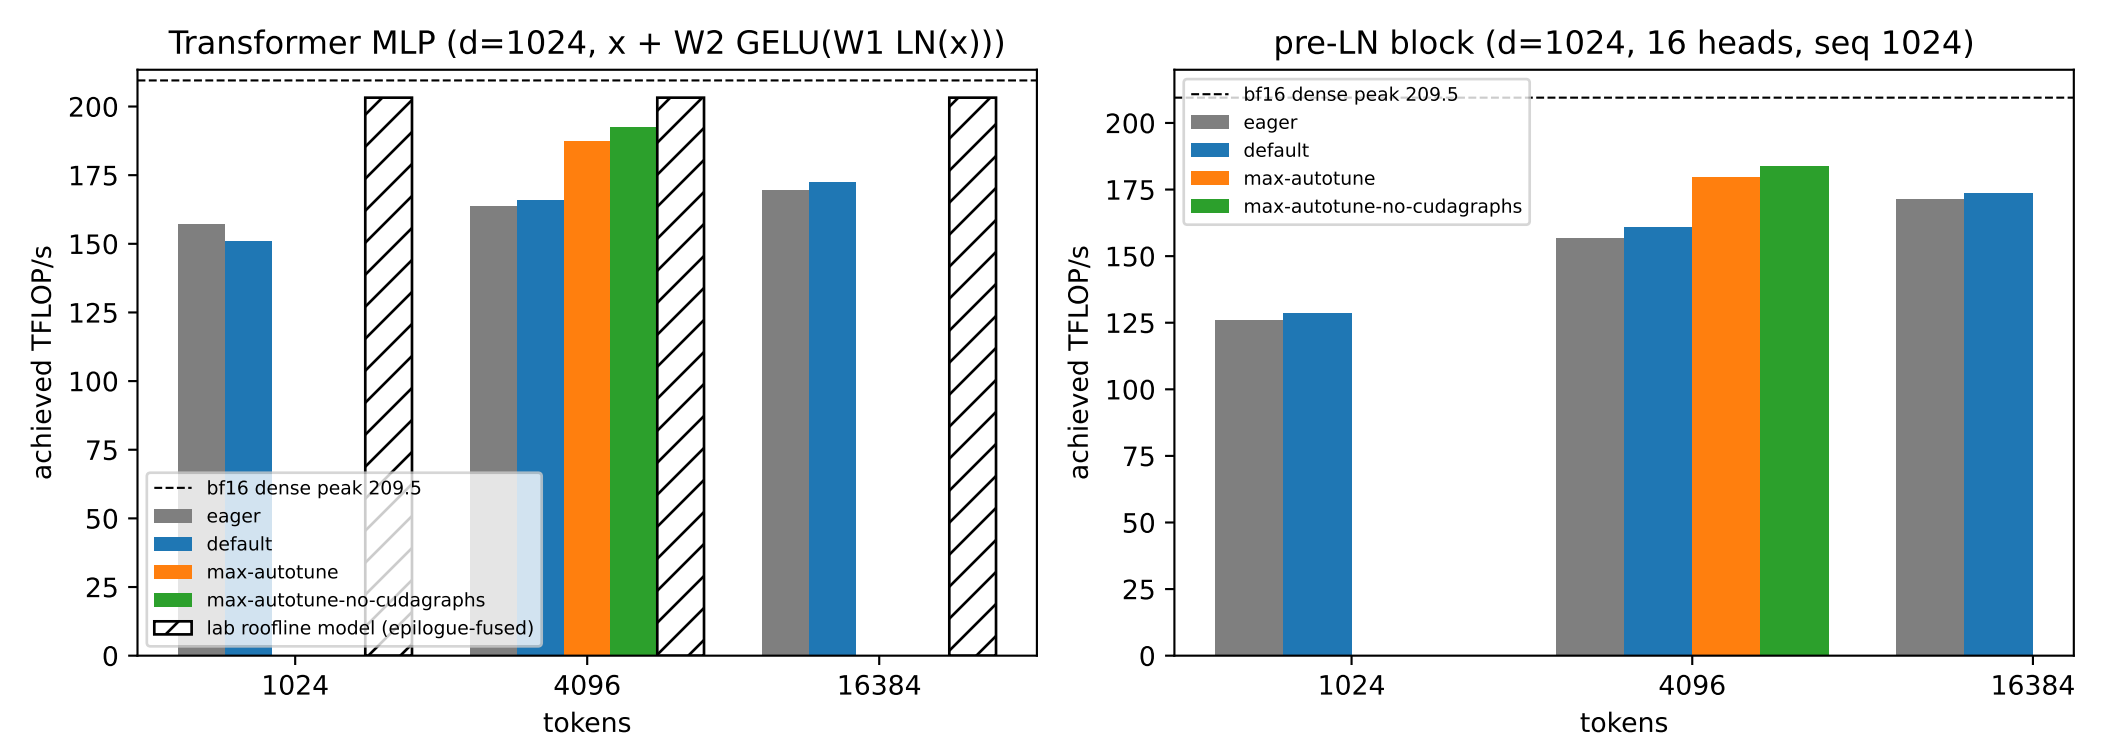
Transformer MLP (d=1024, x + W2 GELU(W1 LN(x)))
achieved TFLOP/s
tokens
1024
4096
16384
0
25
50
75
100
125
150
175
200
bf16 dense peak 209.5
eager
default
max-autotune
max-autotune-no-cudagraphs
lab roofline model (epilogue-fused)
pre-LN block (d=1024, 16 heads, seq 1024)
achieved TFLOP/s
tokens
1024
4096
16384
bf16 dense peak 209.5
eager
default
max-autotune
max-autotune-no-cudagraphs