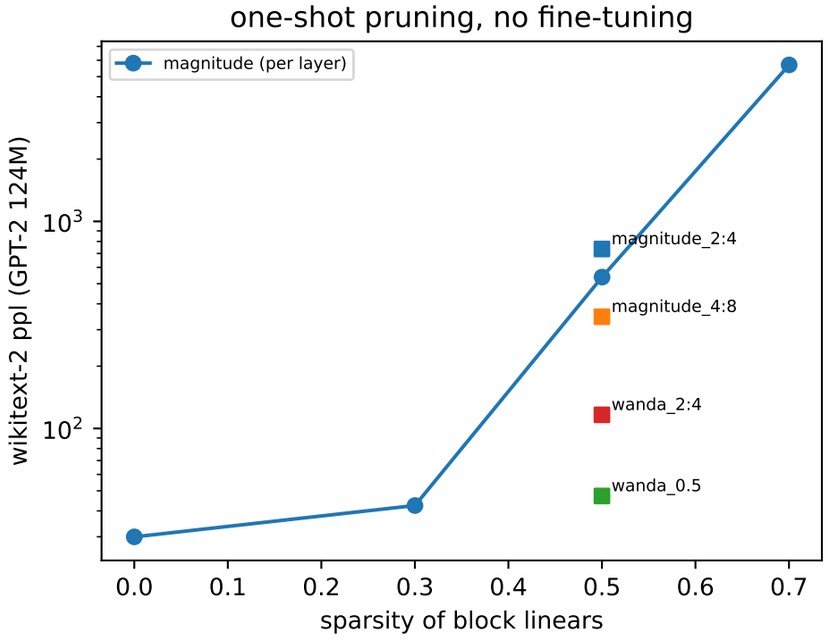
one-shot pruning, no fine-tuning
magnitude (per layer)
magnitude_2:4
magnitude_4:8
wanda_2:4
wanda_0.5
wikitext-2 ppl (GPT-2 124M)
sparsity of block linears

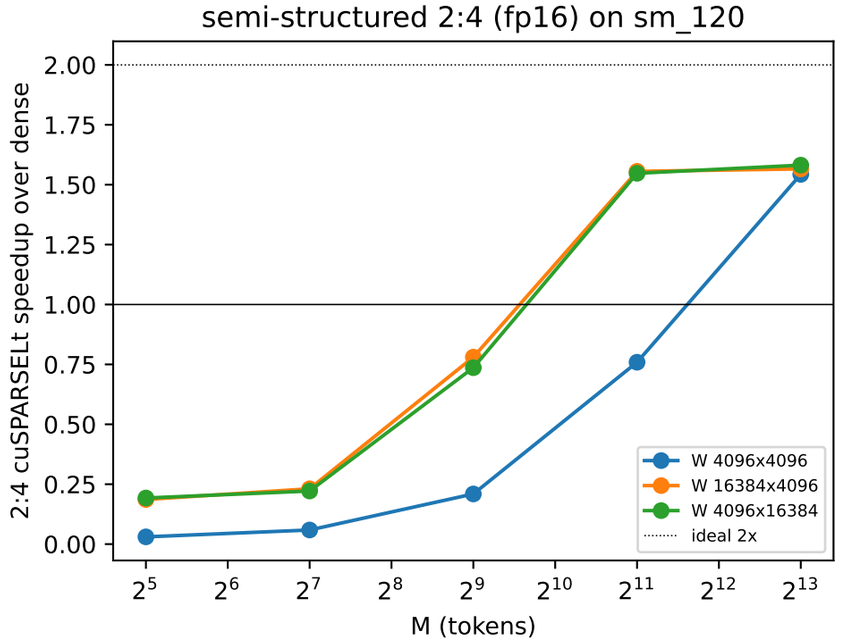
semi-structured 2:4 (fp16) on sm_120
2:4 cuSPARSELt speedup over dense
M (tokens)
W 4096x4096
W 16384x4096
W 4096x16384
ideal 2x

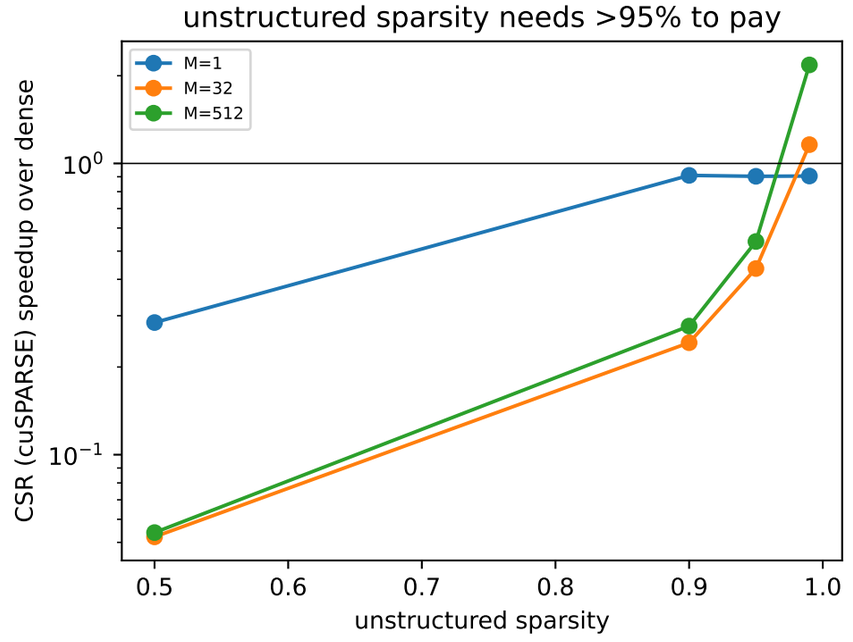
unstructured sparsity needs >95% to pay
M=1
M=32
M=512
CSR (cuSPARSE) speedup over dense
unstructured sparsity