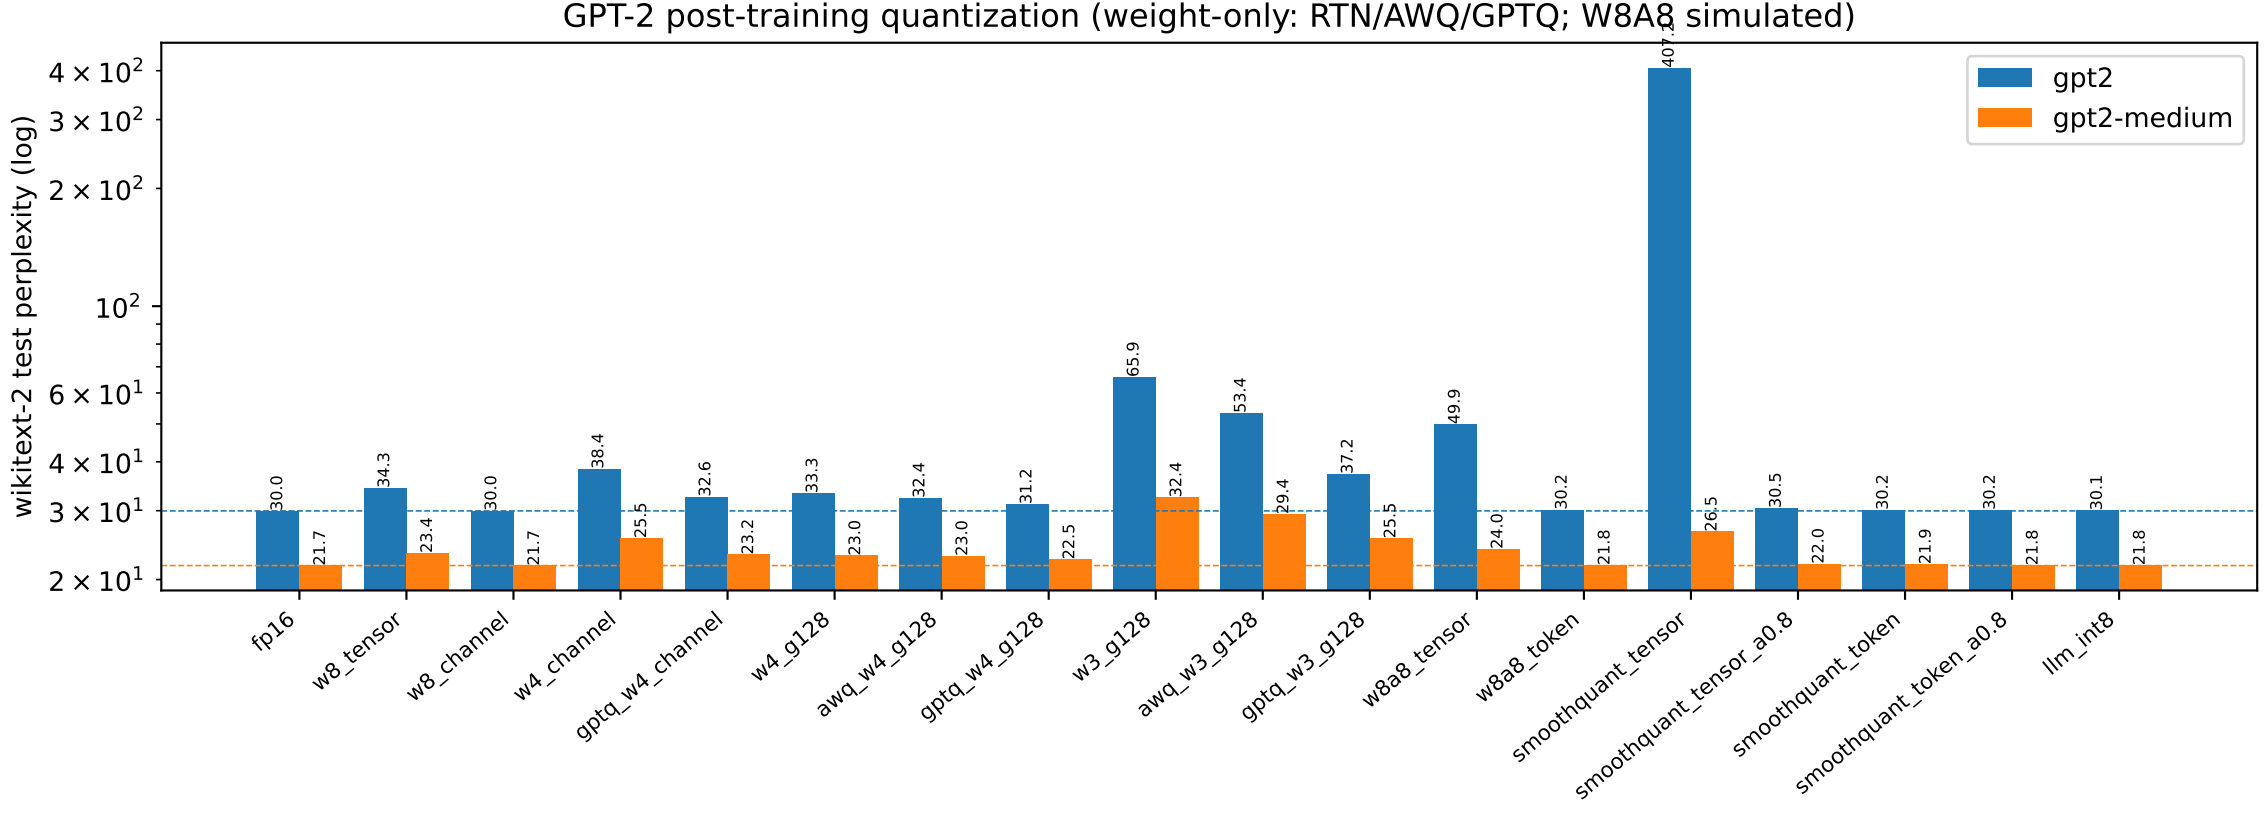
GPT-2 post-training quantization (weight-only: RTN/AWQ/GPTQ; W8A8 simulated)
wikitext-2 test perplexity (log)
gpt2
gpt2-medium
fp16
w8_tensor
w8_channel
w4_channel
gptq_w4_channel
w4_g128
awq_w4_g128
gptq_w4_g128
w3_g128
awq_w3_g128
gptq_w3_g128
w8a8_tensor
w8a8_token
smoothquant_tensor
smoothquant_tensor_a0.8
smoothquant_token
smoothquant_token_a0.8
llm_int8
30.0
21.7
34.3
23.4
30.0
21.7
38.4
25.5
32.6
23.2
33.3
23.0
32.4
23.0
31.2
22.5
65.9
32.4
53.4
29.4
37.2
25.5
49.9
24.0
30.2
21.8
407.2
26.5
30.5
22.0
30.2
21.9
30.2
21.8
30.1
21.8