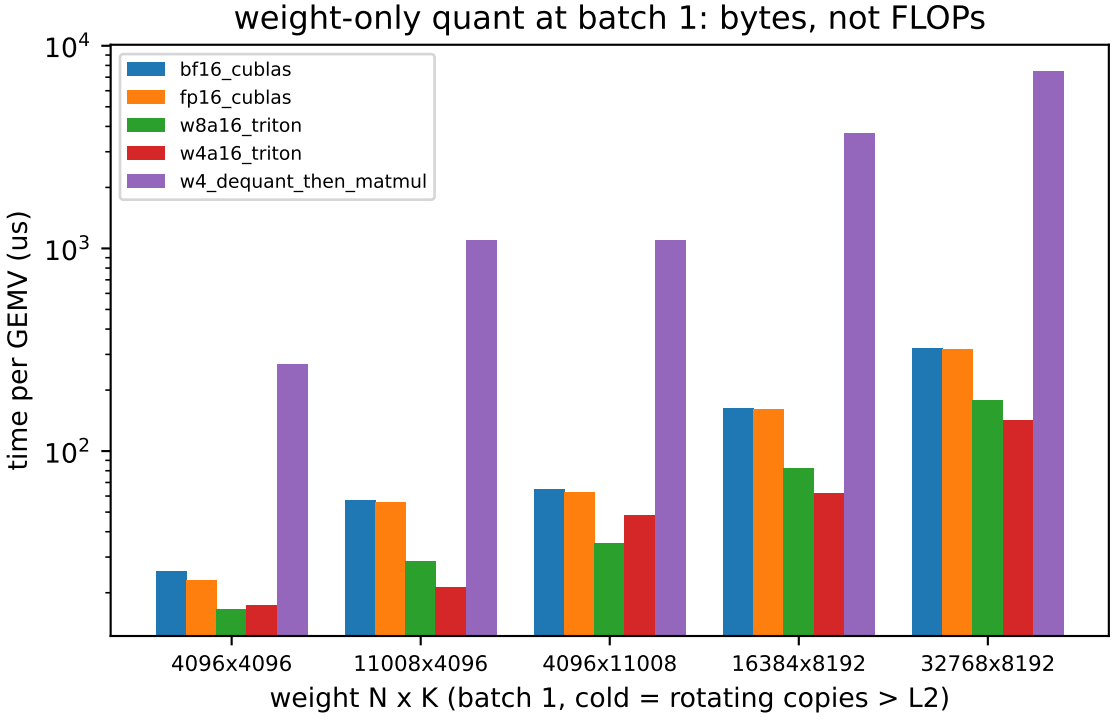
weight-only quant at batch 1: bytes, not FLOPs
time per GEMV (us)
10^4
10^3
10^2
bf16_cublas
fp16_cublas
w8a16_triton
w4a16_triton
w4_dequant_then_matmul
4096x4096
11008x4096
4096x11008
16384x8192
32768x8192
weight N x K (batch 1, cold = rotating copies > L2)

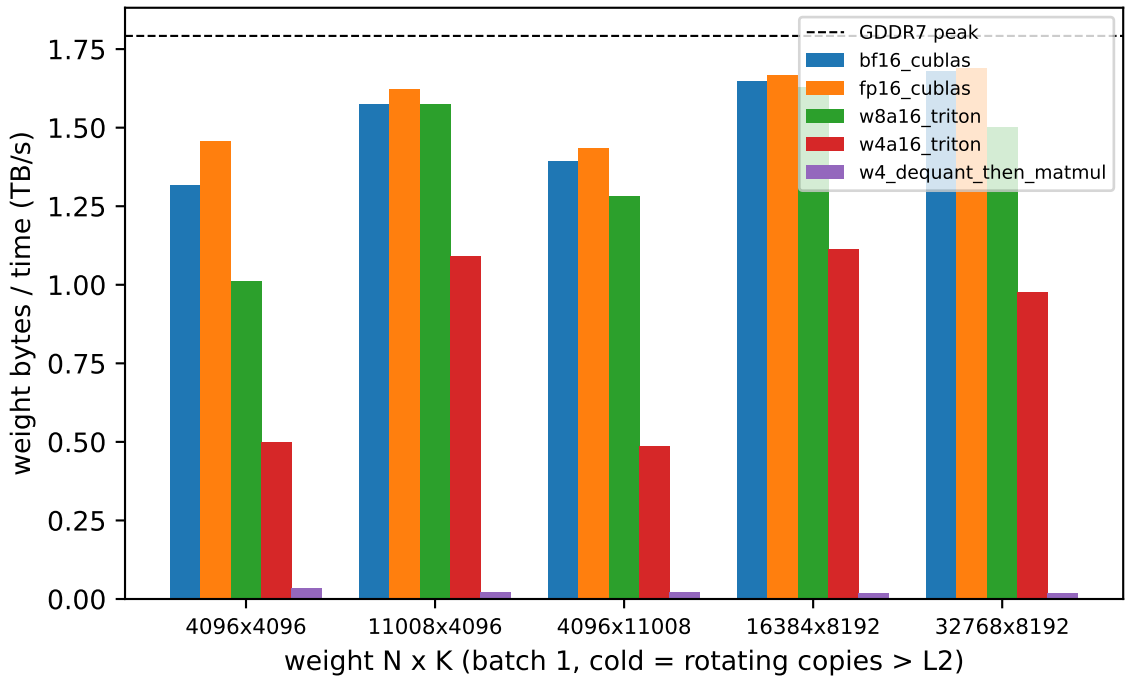
weight bytes / time (TB/s)
1.75
1.50
1.25
1.00
0.75
0.50
0.25
0.00
GDDR7 peak
bf16_cublas
fp16_cublas
w8a16_triton
w4a16_triton
w4_dequant_then_matmul
4096x4096
11008x4096
4096x11008
16384x8192
32768x8192
weight N x K (batch 1, cold = rotating copies > L2)