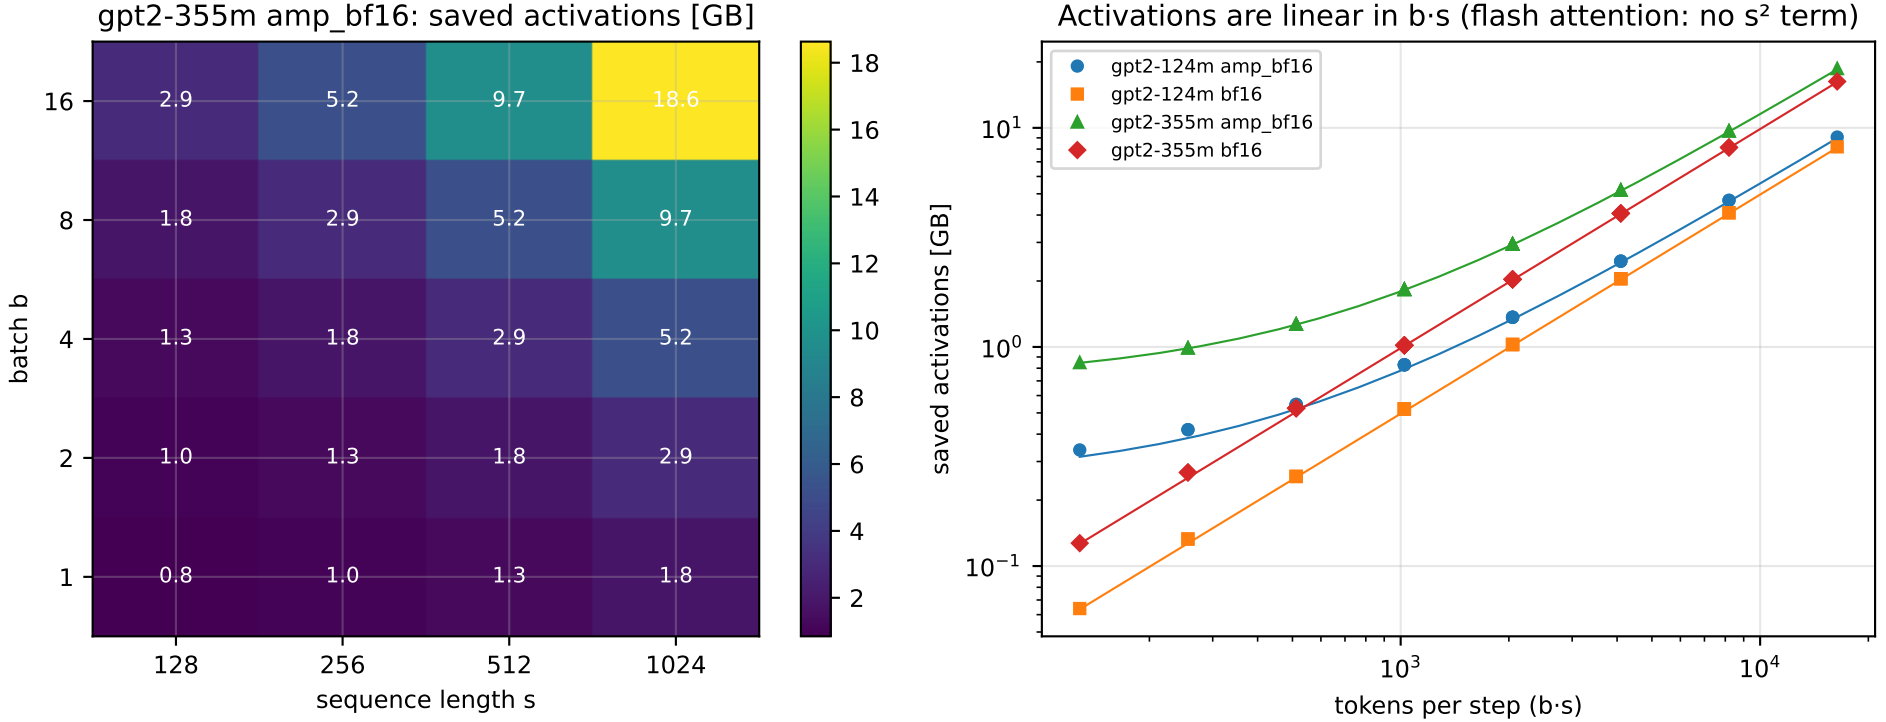
gpt2-355m amp_bf16: saved activations [GB]
batch b
sequence length s
128
256
512
1024
16
8
4
2
1
2.9
5.2
9.7
18.6
1.8
2.9
5.2
9.7
1.3
1.8
2.9
5.2
1.0
1.3
1.8
2.9
0.8
1.0
1.3
1.8
18
16
14
12
10
8
6
4
2
Activations are linear in b·s (flash attention: no s² term)
saved activations [GB]
tokens per step (b·s)
gpt2-124m amp_bf16
gpt2-124m bf16
gpt2-355m amp_bf16
gpt2-355m bf16
10^{1}
10^{0}
10^{-1}
10^{3}
10^{4}