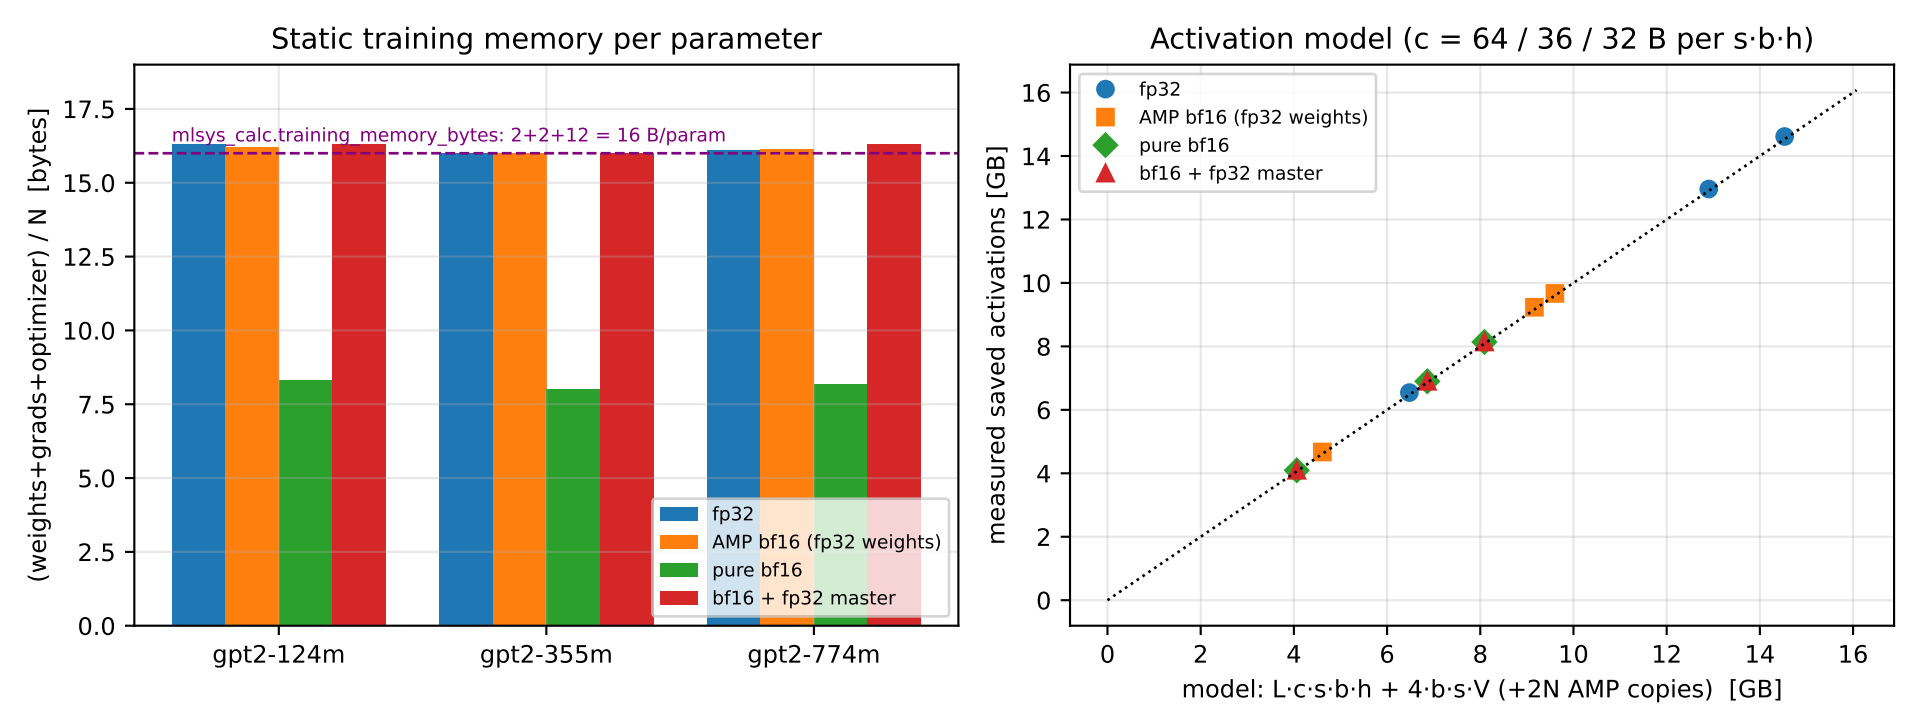
Static training memory per parameter
(weights+grads+optimizer) / N [bytes]
mlsys_calc.training_memory_bytes: 2+2+12 = 16 B/param
fp32
AMP bf16 (fp32 weights)
pure bf16
bf16 + fp32 master
gpt2-124m
gpt2-355m
gpt2-774m
Activation model (c = 64 / 36 / 32 B per s·b·h)
measured saved activations [GB]
model: L·c·s·b·h + 4·b·s·V (+2N AMP copies) [GB]