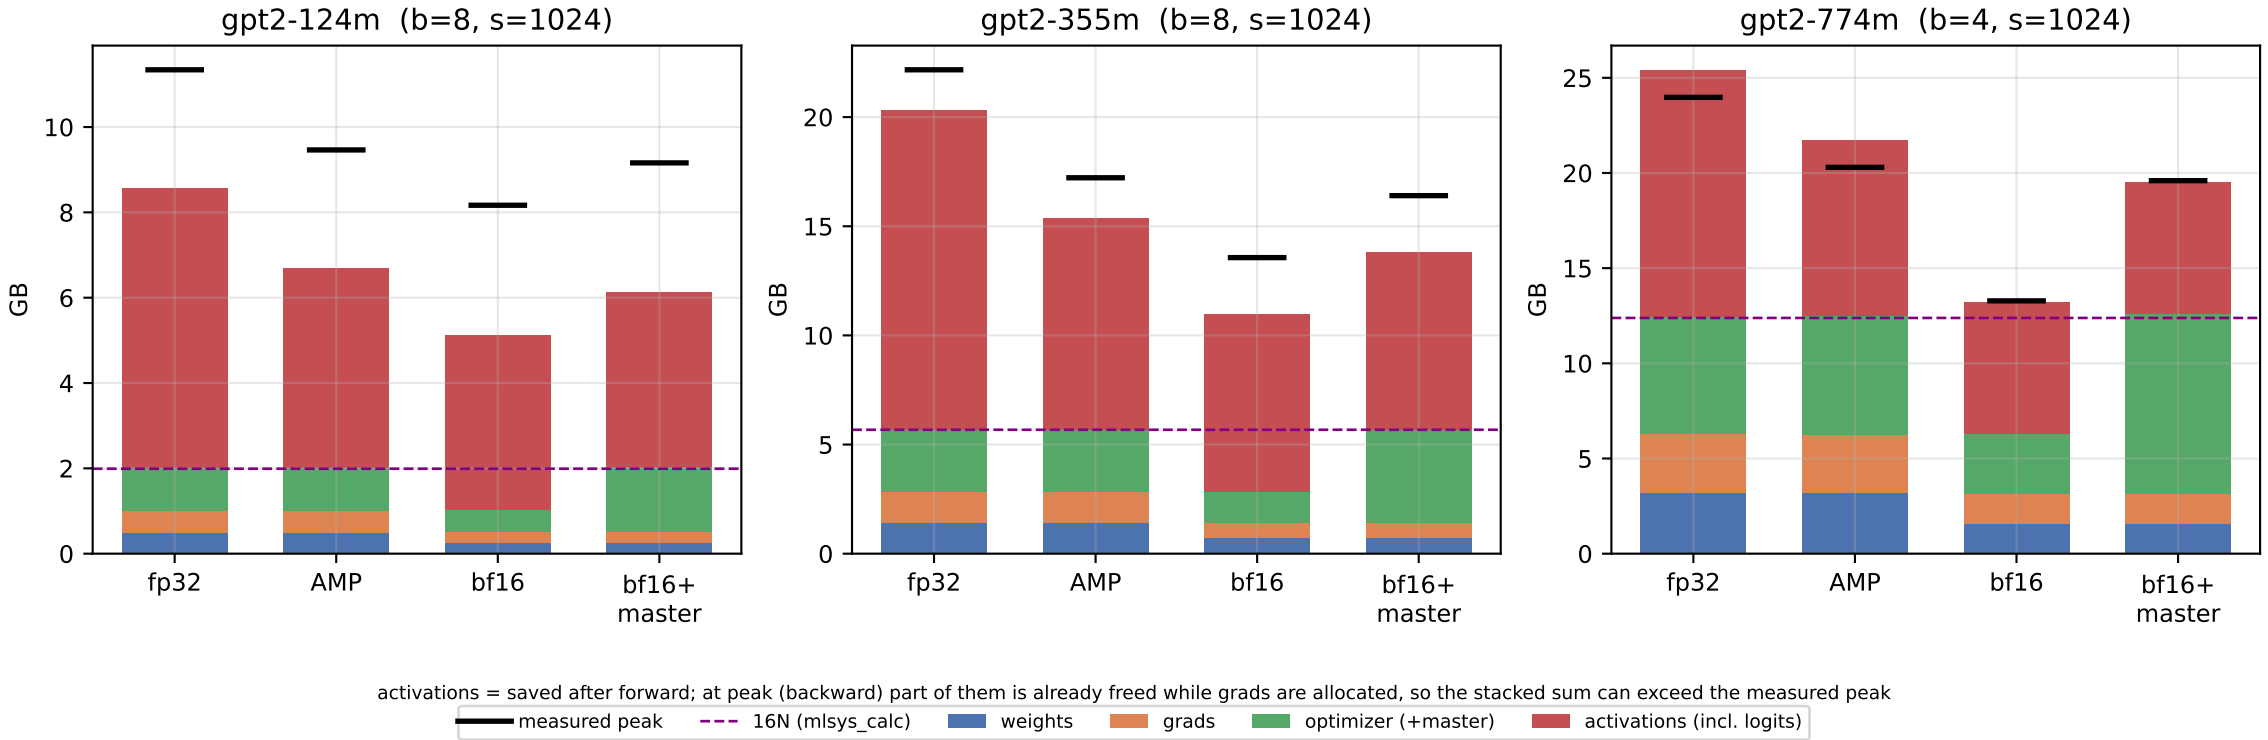
gpt2-124m (b=8, s=1024)
gpt2-355m (b=8, s=1024)
gpt2-774m (b=4, s=1024)
GB
fp32
AMP
bf16
bf16+ master
activations = saved after forward; at peak (backward) part of them is already freed while grads are allocated, so the stacked sum can exceed the measured peak
measured peak
16N (mlsys_calc)
weights
grads
optimizer (+master)
activations (incl. logits)