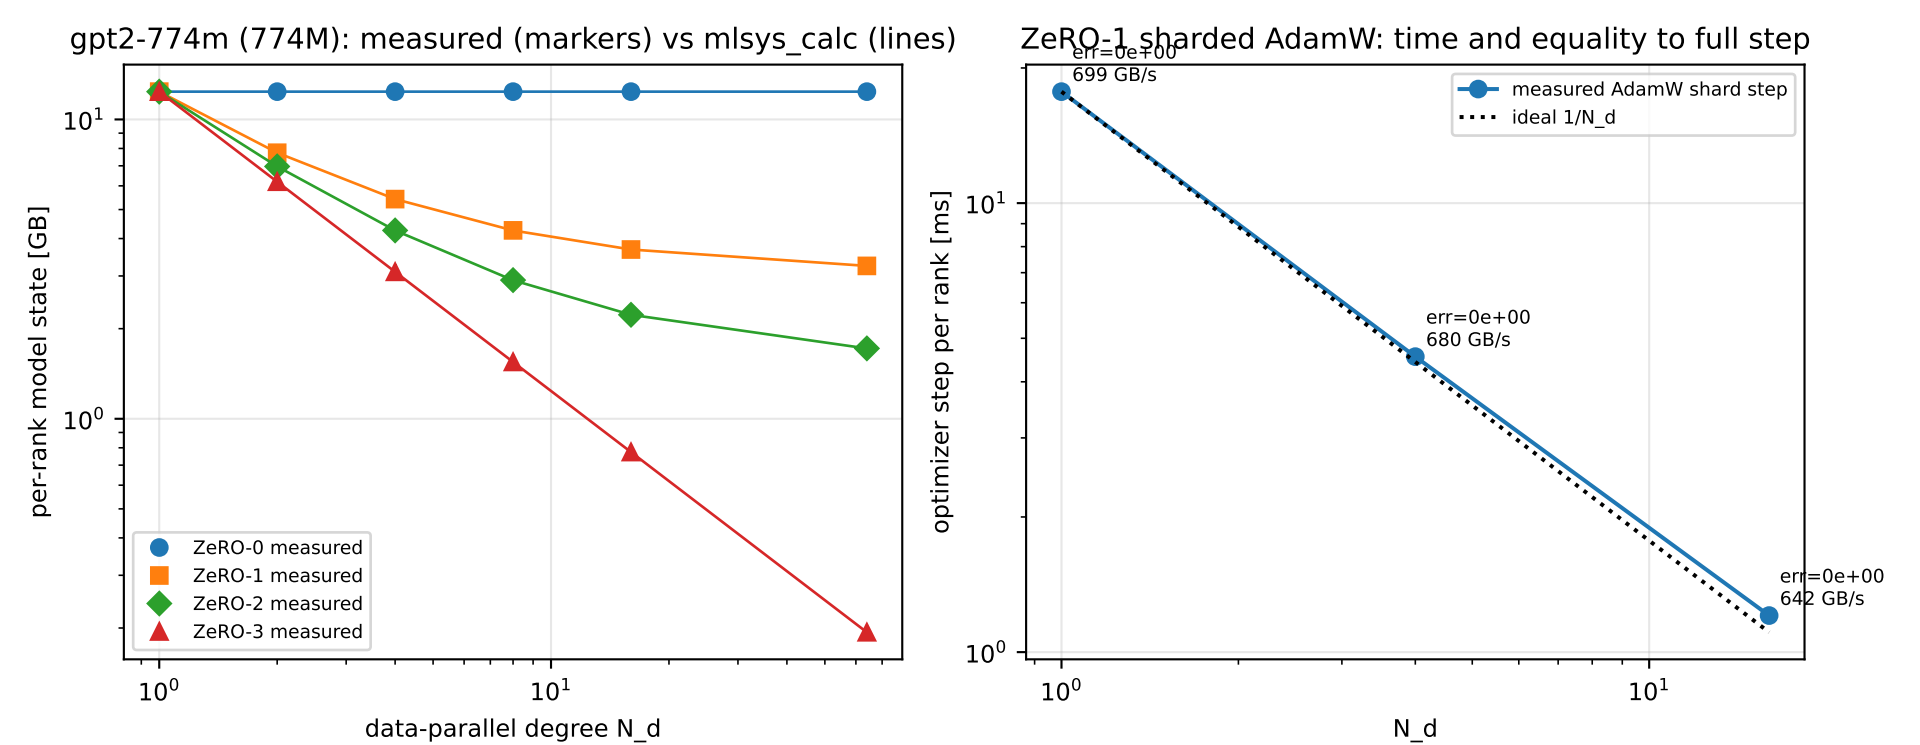
gpt2-774m (774M): measured (markers) vs mlsys_calc (lines)
per-rank model state [GB]
data-parallel degree N_d
ZeRO-0 measured
ZeRO-1 measured
ZeRO-2 measured
ZeRO-3 measured
ZeRO-1 sharded AdamW: time and equality to full step
optimizer step per rank [ms]
N_d
measured AdamW shard step
ideal 1/N_d
err=0e+00
699 GB/s
err=0e+00
680 GB/s
err=0e+00
642 GB/s