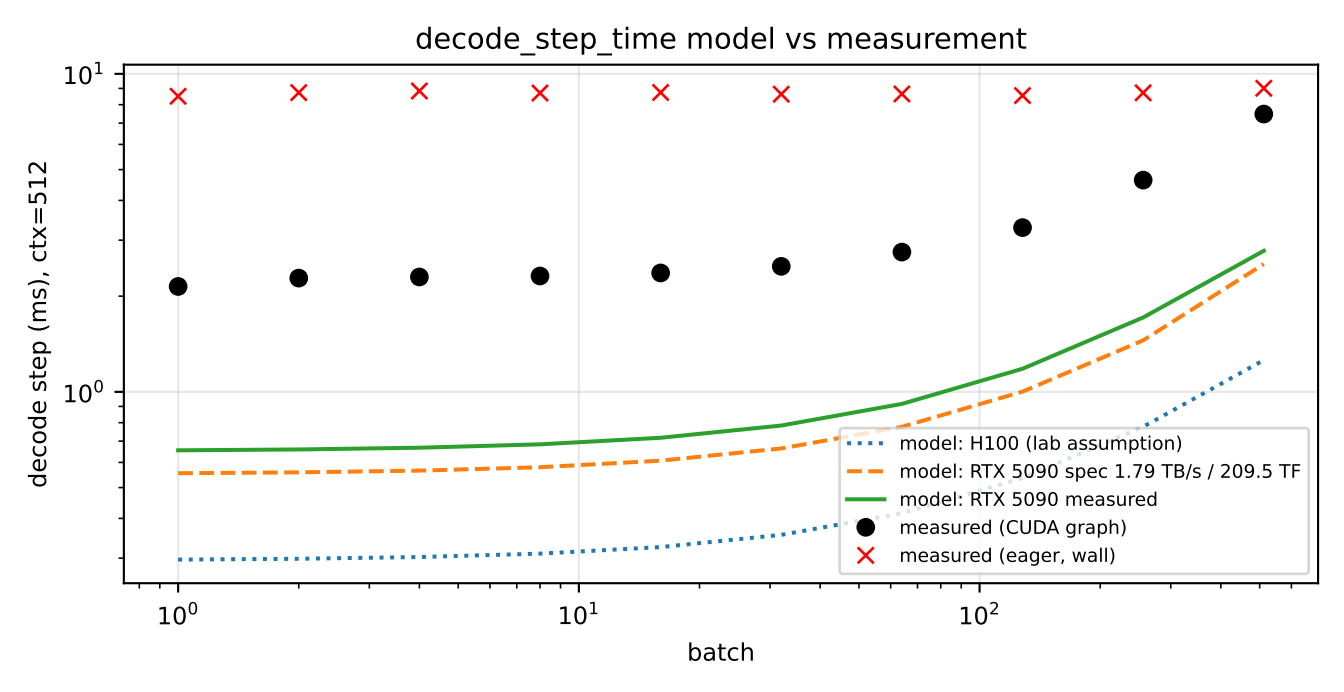
decode_step_time model vs measurement
decode step (ms), ctx=512
batch
10^1
10^0
10^0
10^1
10^2
model: H100 (lab assumption)
model: RTX 5090 spec 1.79 TB/s / 209.5 TF
model: RTX 5090 measured
measured (CUDA graph)
measured (eager, wall)