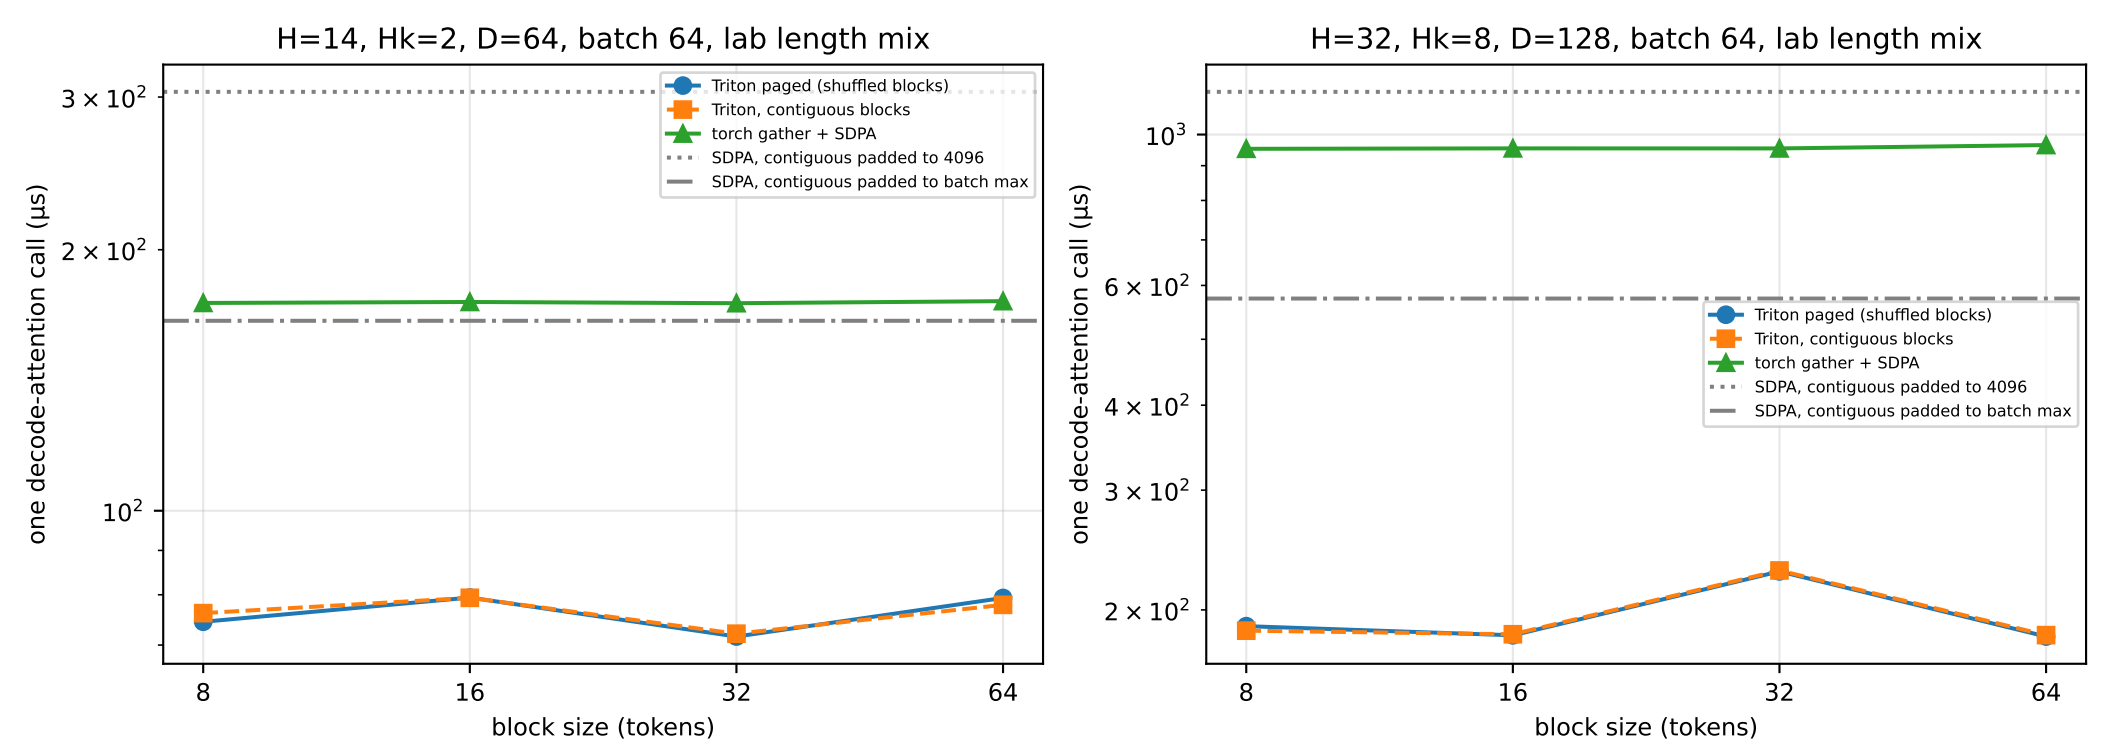
H=14, Hk=2, D=64, batch 64, lab length mix
one decode-attention call (µs)
block size (tokens)
Triton paged (shuffled blocks)
Triton, contiguous blocks
torch gather + SDPA
SDPA, contiguous padded to 4096
SDPA, contiguous padded to batch max
H=32, Hk=8, D=128, batch 64, lab length mix
one decode-attention call (µs)
block size (tokens)
Triton paged (shuffled blocks)
Triton, contiguous blocks
torch gather + SDPA
SDPA, contiguous padded to 4096
SDPA, contiguous padded to batch max