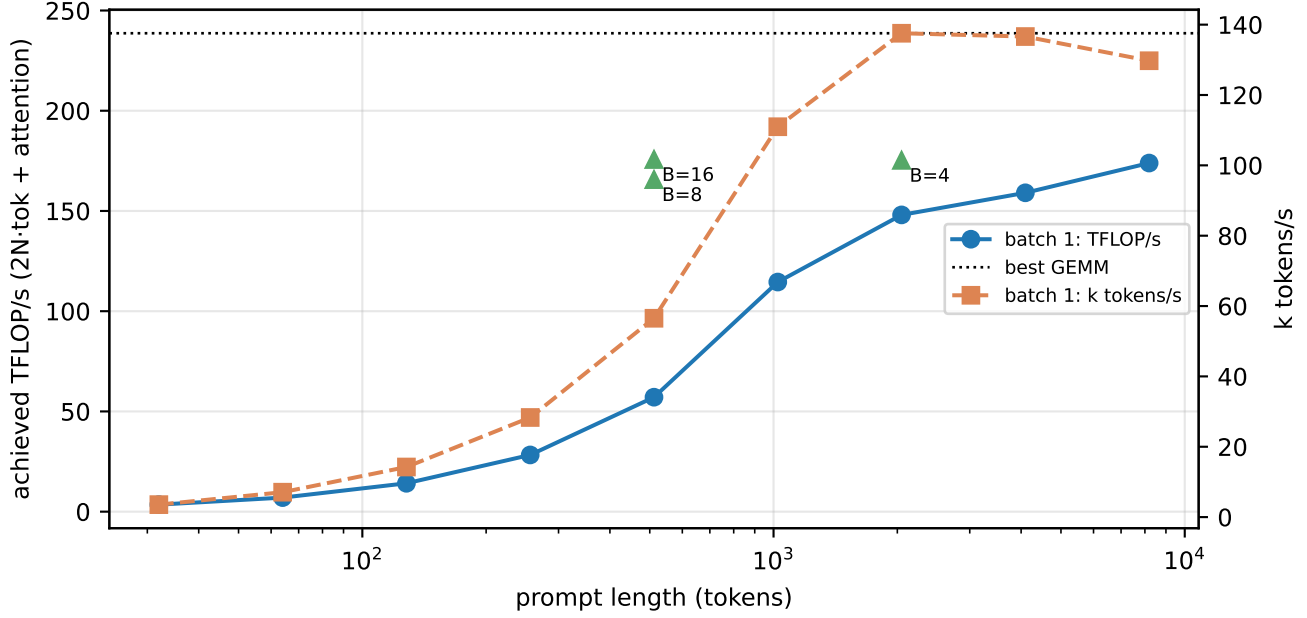
Prefill is compute-bound once tokens ≳ ridge
achieved TFLOP/s (2N·tok + attention)
k tokens/s
prompt length (tokens)
250
200
150
100
50
0
140
120
100
80
60
40
20
0
10^2
10^3
10^4
B=16
B=8
B=4
batch 1: TFLOP/s
best GEMM
batch 1: k tokens/s