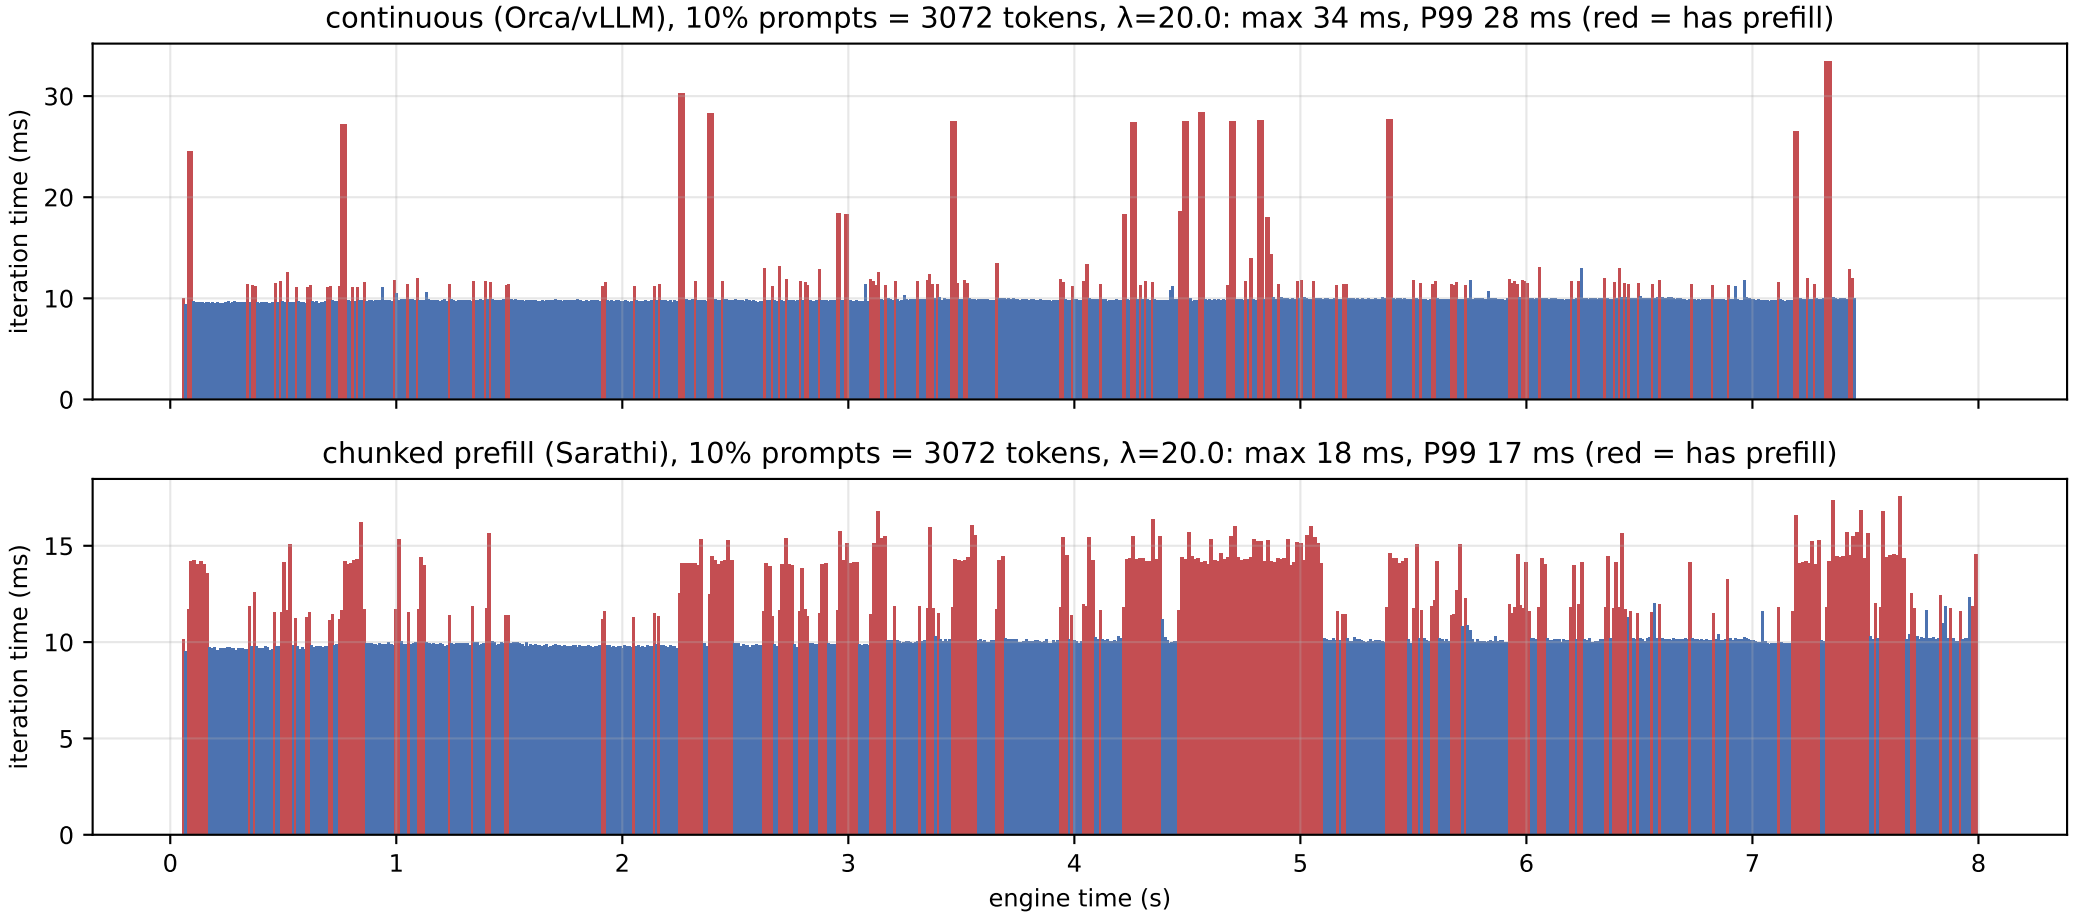
continuous (Orca/vLLM), 10% prompts = 3072 tokens, λ=20.0: max 34 ms, P99 28 ms (red = has prefill)
iteration time (ms)
0
10
20
30
chunked prefill (Sarathi), 10% prompts = 3072 tokens, λ=20.0: max 18 ms, P99 17 ms (red = has prefill)
iteration time (ms)
0
5
10
15
0
1
2
3
4
5
6
7
8
engine time (s)