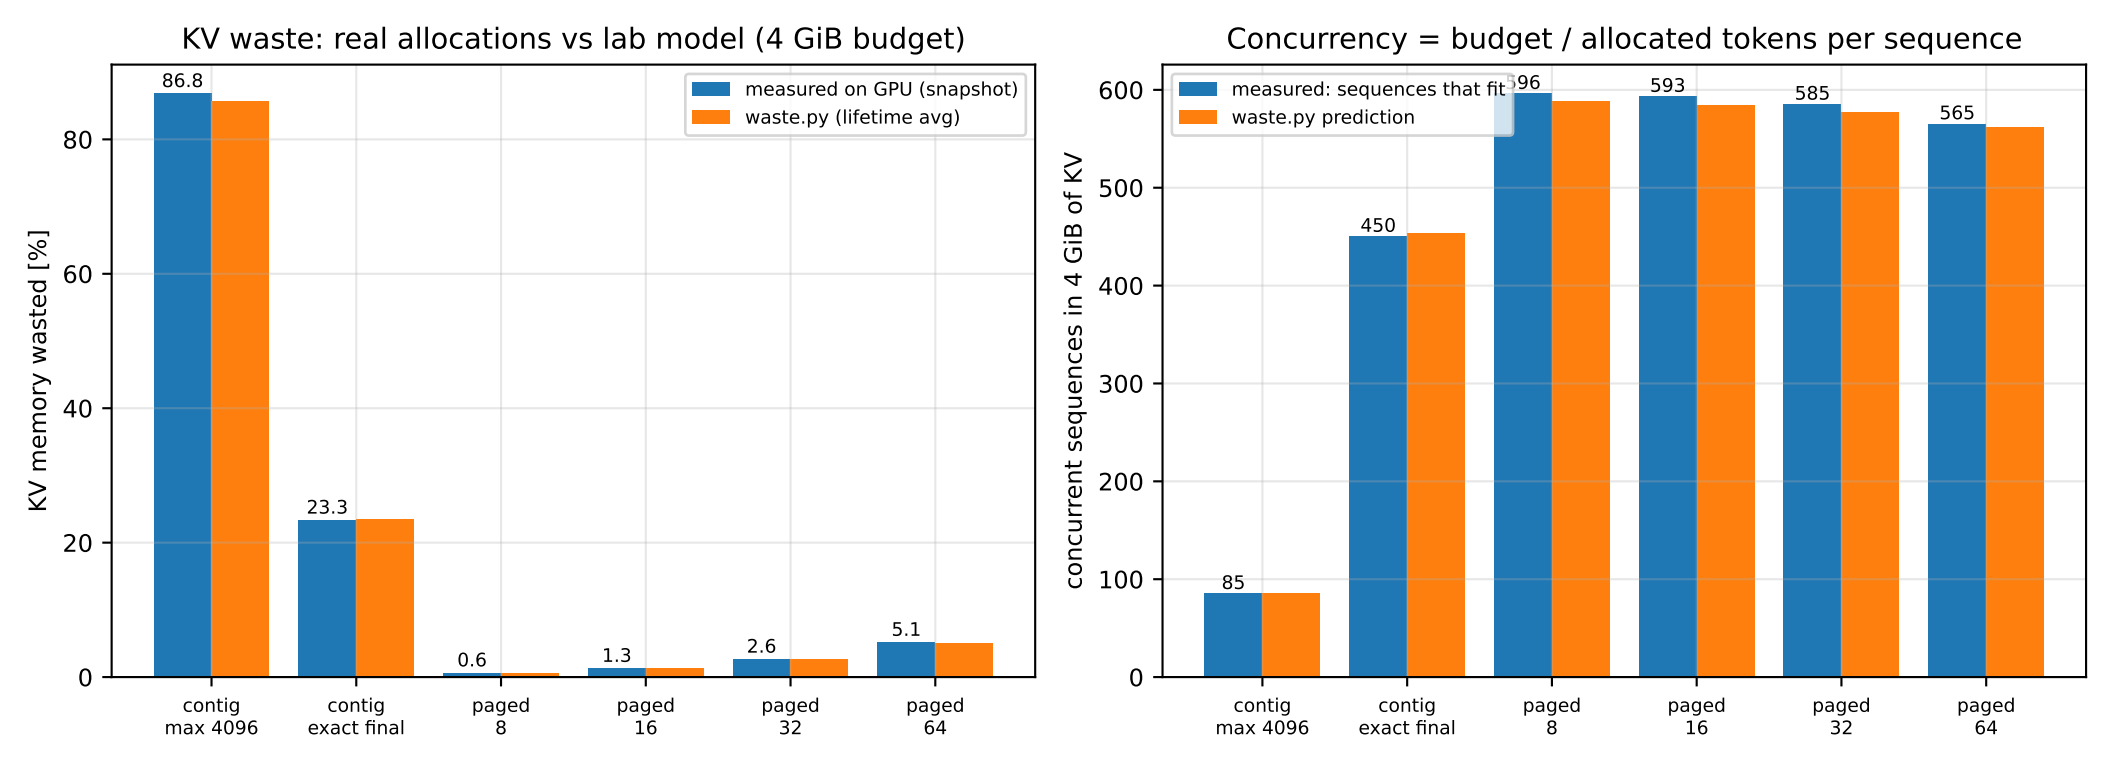
KV waste: real allocations vs lab model (4 GiB budget)
KV memory wasted [%]
measured on GPU (snapshot)
waste.py (lifetime avg)
86.8
23.3
0.6
1.3
2.6
5.1
0
20
40
60
80
contig max 4096
contig exact final
paged 8
paged 16
paged 32
paged 64
Concurrency = budget / allocated tokens per sequence
concurrent sequences in 4 GiB of KV
measured: sequences that fit
waste.py prediction
85
450
596
593
585
565
0
100
200
300
400
500
600
contig max 4096
contig exact final
paged 8
paged 16
paged 32
paged 64