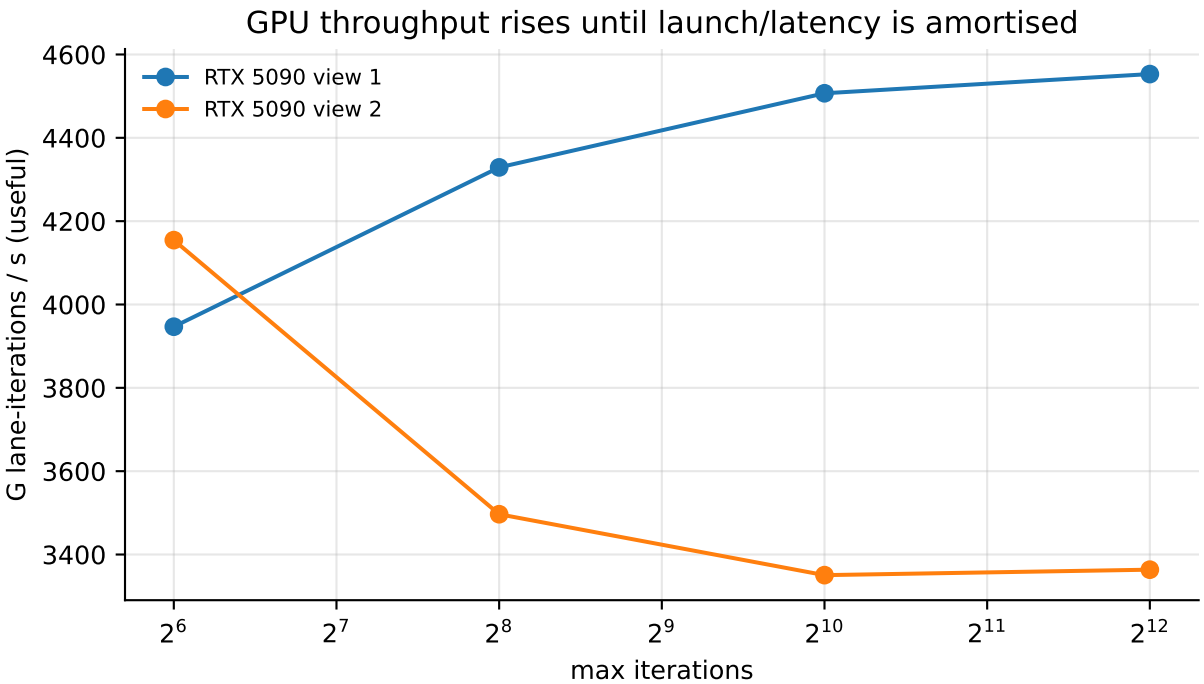
GPU throughput rises until launch/latency is amortised
RTX 5090 view 1
RTX 5090 view 2
G lane-iterations / s (useful)
4600
4400
4200
4000
3800
3600
3400
2^6
2^7
2^8
2^9
2^{10}
2^{11}
2^{12}
max iterations