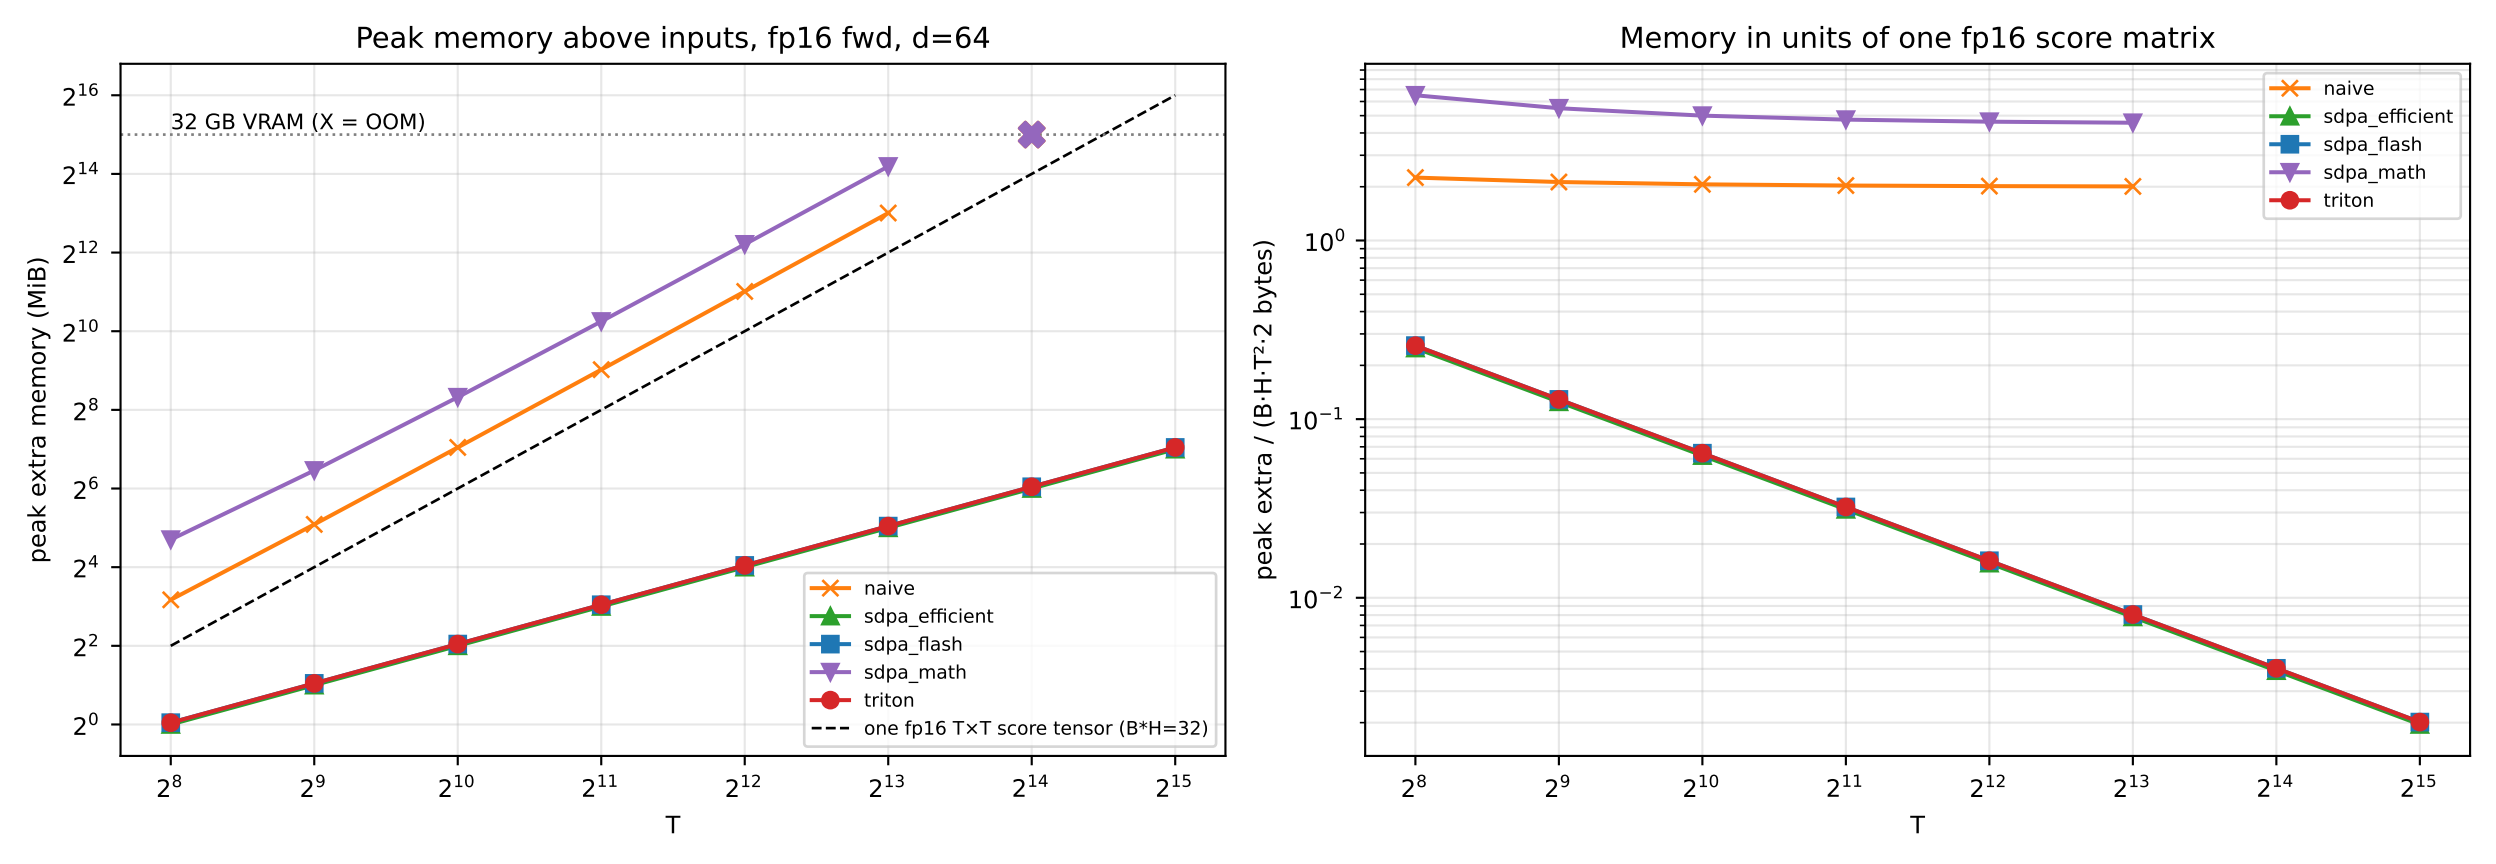
Peak memory above inputs, fp16 fwd, d=64
32 GB VRAM (X = OOM)
peak extra memory (MiB)
T
naive
sdpa_efficient
sdpa_flash
sdpa_math
triton
one fp16 T×T score tensor (B*H=32)
Memory in units of one fp16 score matrix
peak extra / (B·H·T²·2 bytes)
T
naive
sdpa_efficient
sdpa_flash
sdpa_math
triton