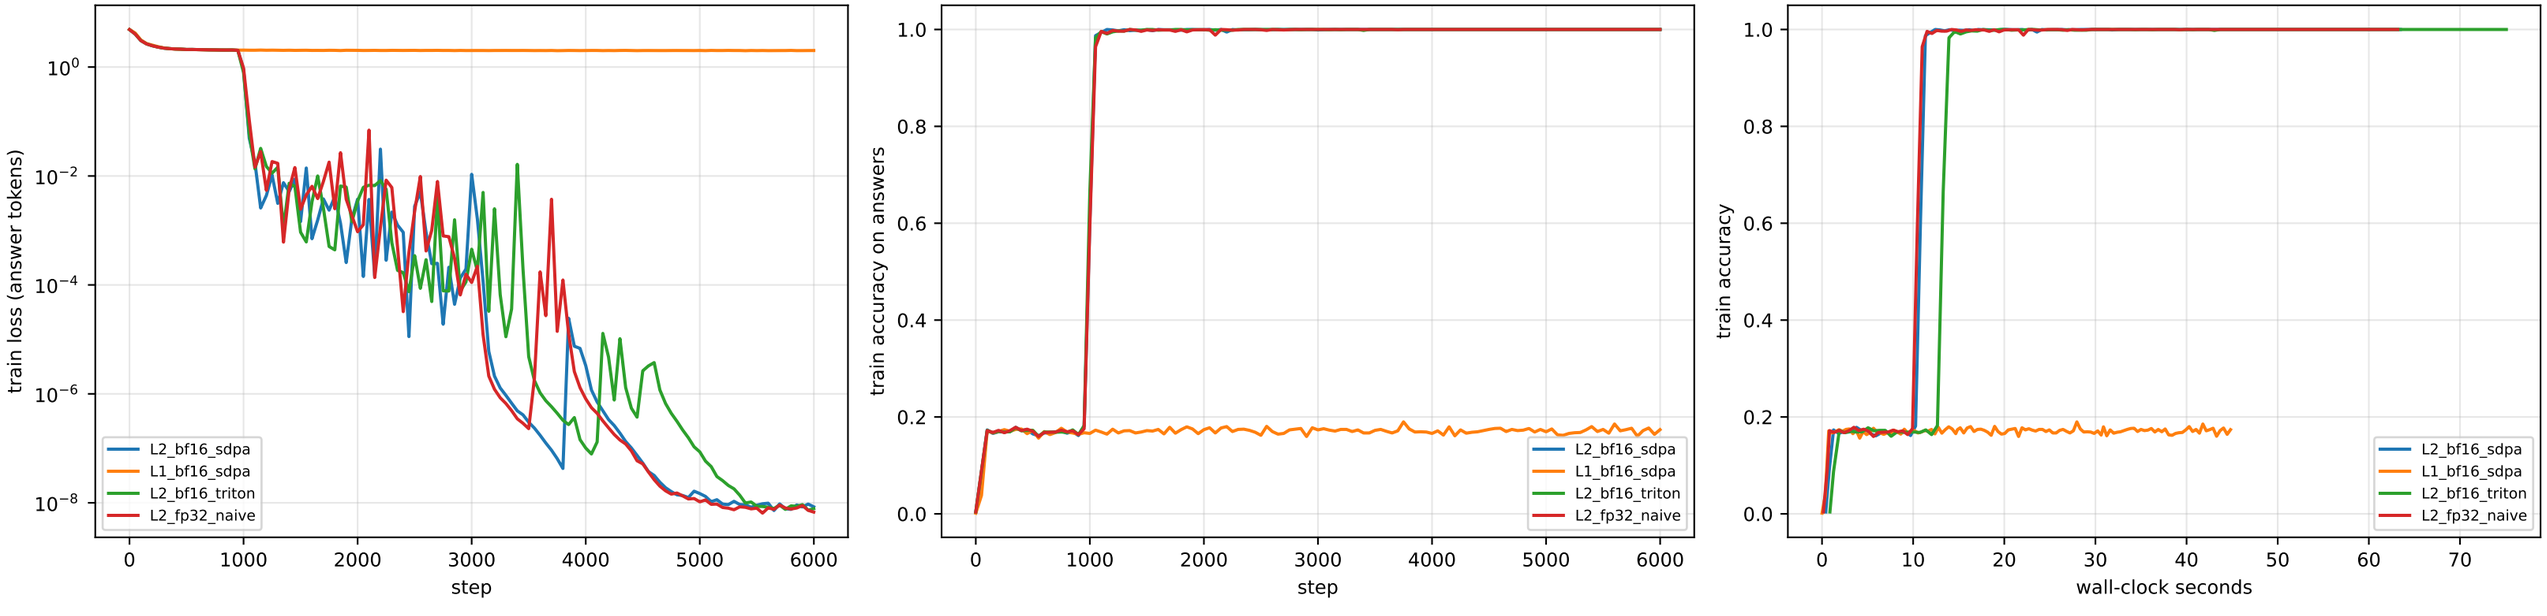
MQAR learning curves (batch 256)
train loss (answer tokens)
step
L2_bf16_sdpa
L1_bf16_sdpa
L2_bf16_triton
L2_fp32_naive
Accuracy vs step
train accuracy on answers
step
Accuracy vs wall clock
train accuracy
wall-clock seconds