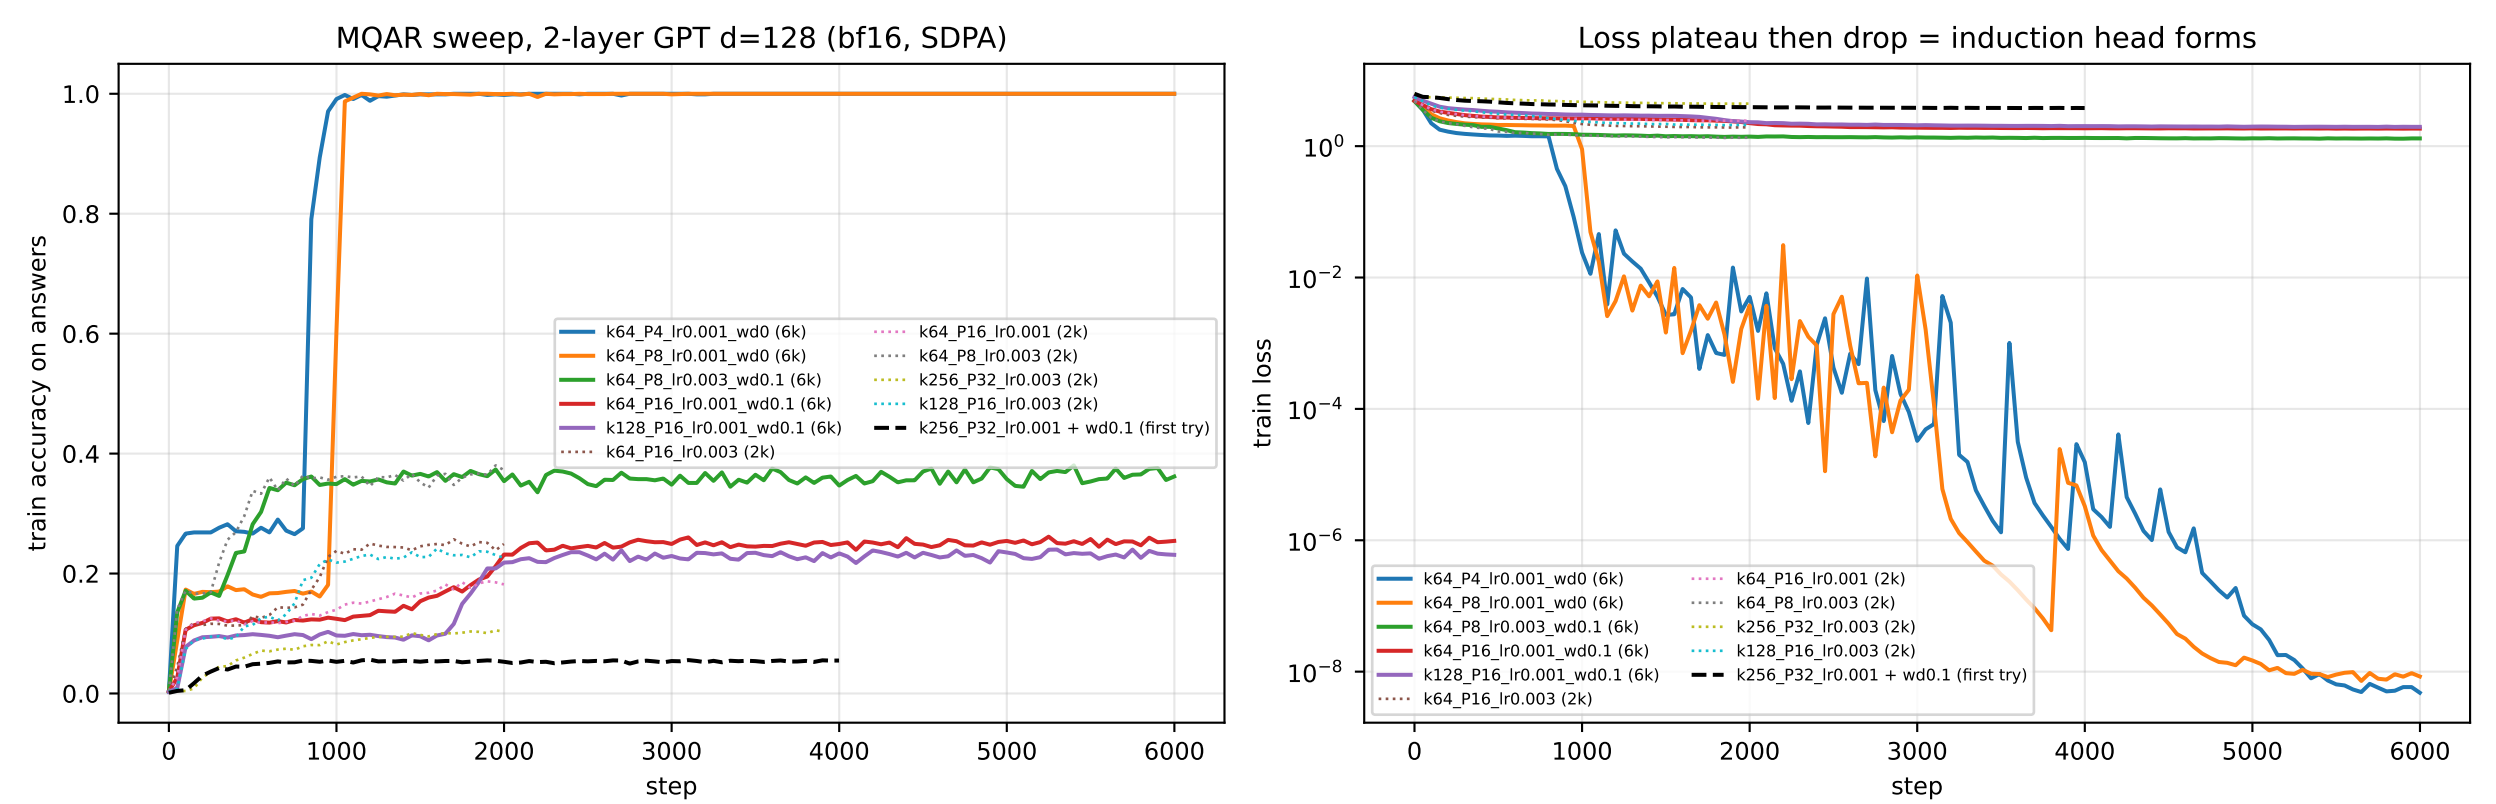
MQAR sweep, 2-layer GPT d=128 (bf16, SDPA)
train accuracy on answers
step
k64_P4_lr0.001_wd0 (6k)
k64_P8_lr0.001_wd0 (6k)
k64_P8_lr0.003_wd0.1 (6k)
k64_P16_lr0.001_wd0.1 (6k)
k128_P16_lr0.001_wd0.1 (6k)
k64_P16_lr0.003 (2k)
k64_P16_lr0.001 (2k)
k64_P8_lr0.003 (2k)
k256_P32_lr0.003 (2k)
k128_P16_lr0.003 (2k)
k256_P32_lr0.001 + wd0.1 (first try)
Loss plateau then drop = induction head forms
train loss
step