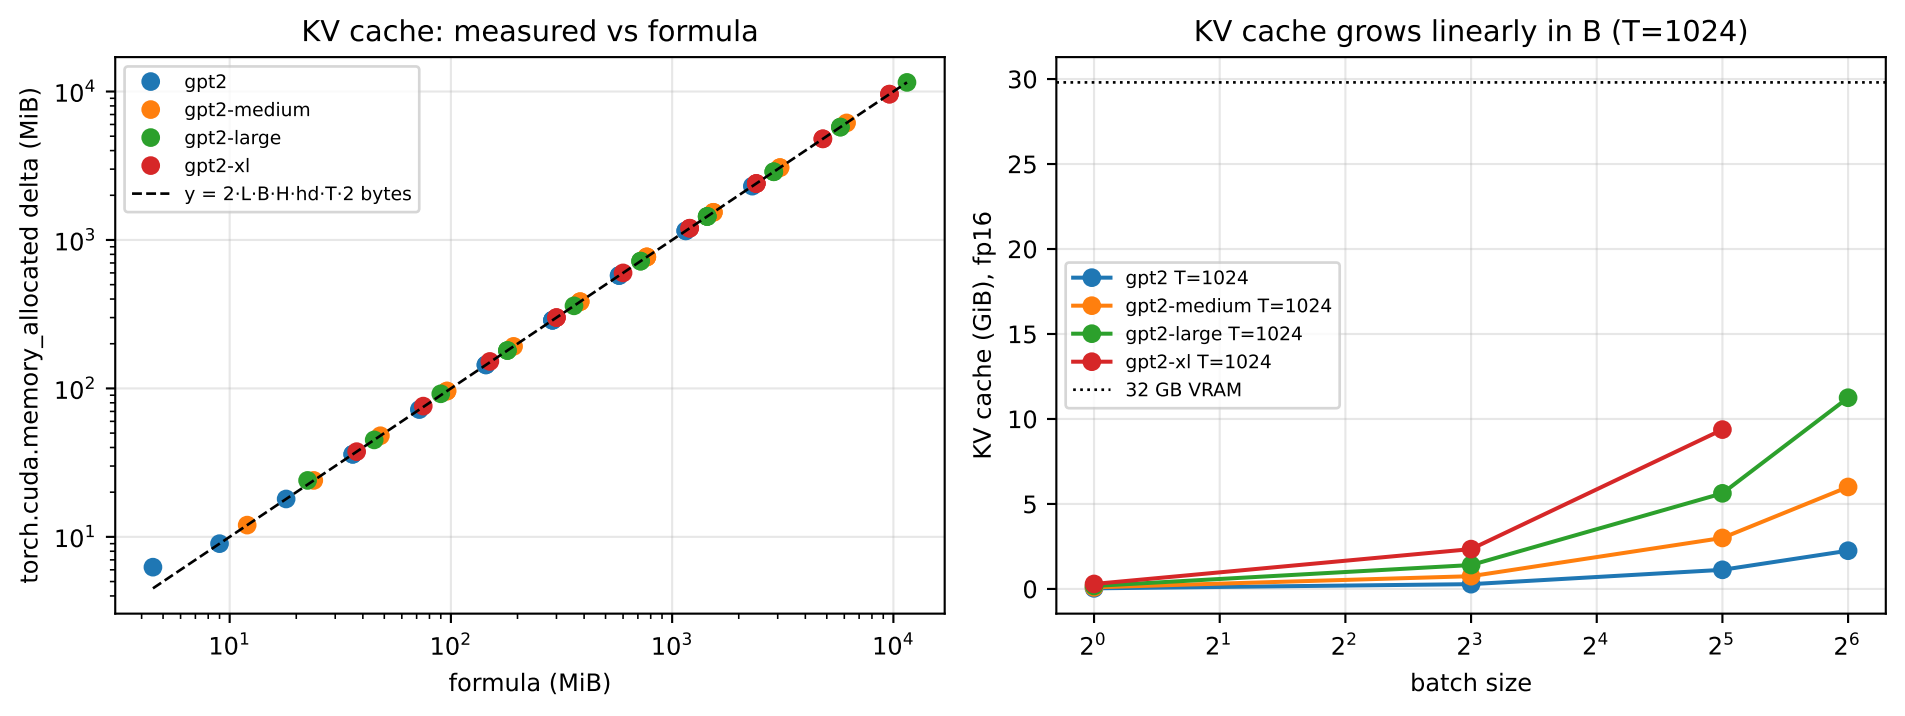
KV cache: measured vs formula
torch.cuda.memory_allocated delta (MiB)
formula (MiB)
gpt2
gpt2-medium
gpt2-large
gpt2-xl
y = 2·L·B·H·hd·T·2 bytes
KV cache grows linearly in B (T=1024)
KV cache (GiB), fp16
batch size
gpt2 T=1024
gpt2-medium T=1024
gpt2-large T=1024
gpt2-xl T=1024
32 GB VRAM