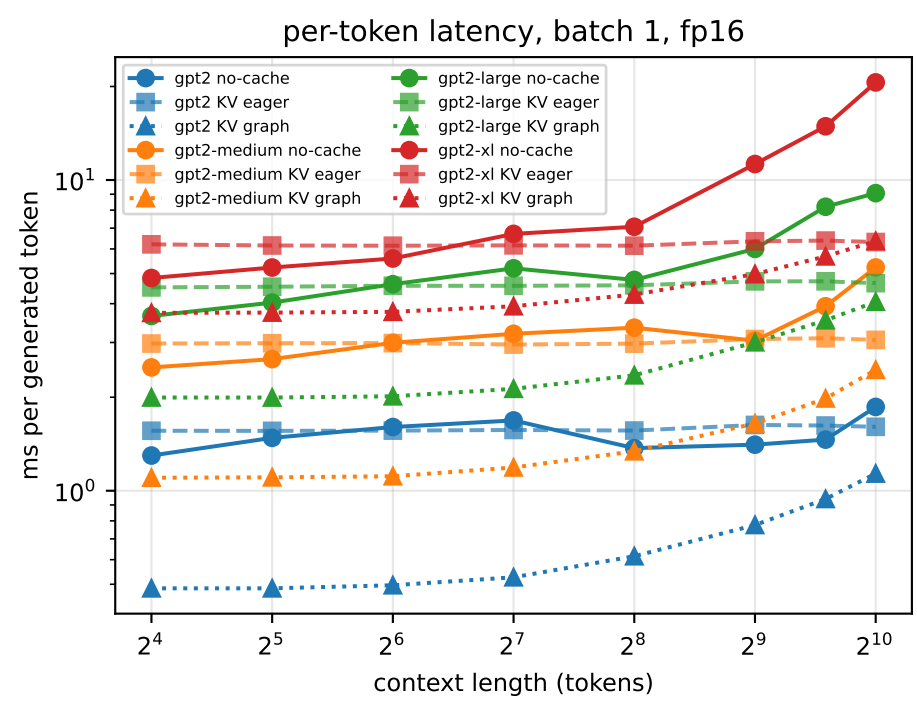
per-token latency, batch 1, fp16
ms per generated token
context length (tokens)
gpt2 no-cache
gpt2 KV eager
gpt2 KV graph
gpt2-medium no-cache
gpt2-medium KV eager
gpt2-medium KV graph
gpt2-large no-cache
gpt2-large KV eager
gpt2-large KV graph
gpt2-xl no-cache
gpt2-xl KV eager
gpt2-xl KV graph

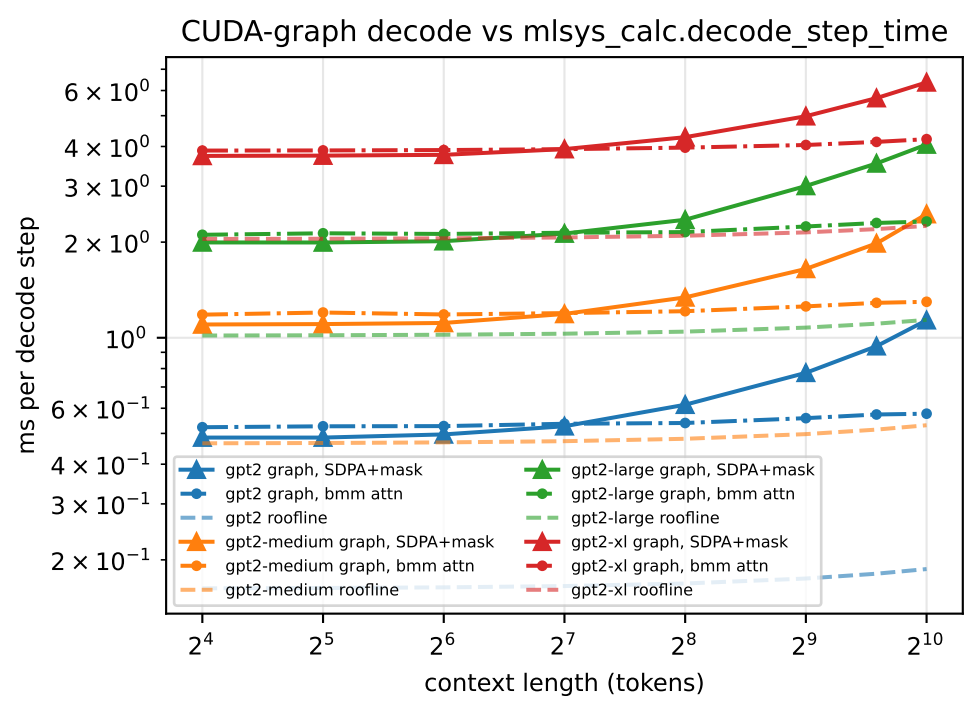
CUDA-graph decode vs mlsys_calc.decode_step_time
ms per decode step
context length (tokens)
gpt2 graph, SDPA+mask
gpt2 graph, bmm attn
gpt2 roofline
gpt2-medium graph, SDPA+mask
gpt2-medium graph, bmm attn
gpt2-medium roofline
gpt2-large graph, SDPA+mask
gpt2-large graph, bmm attn
gpt2-large roofline
gpt2-xl graph, SDPA+mask
gpt2-xl graph, bmm attn
gpt2-xl roofline