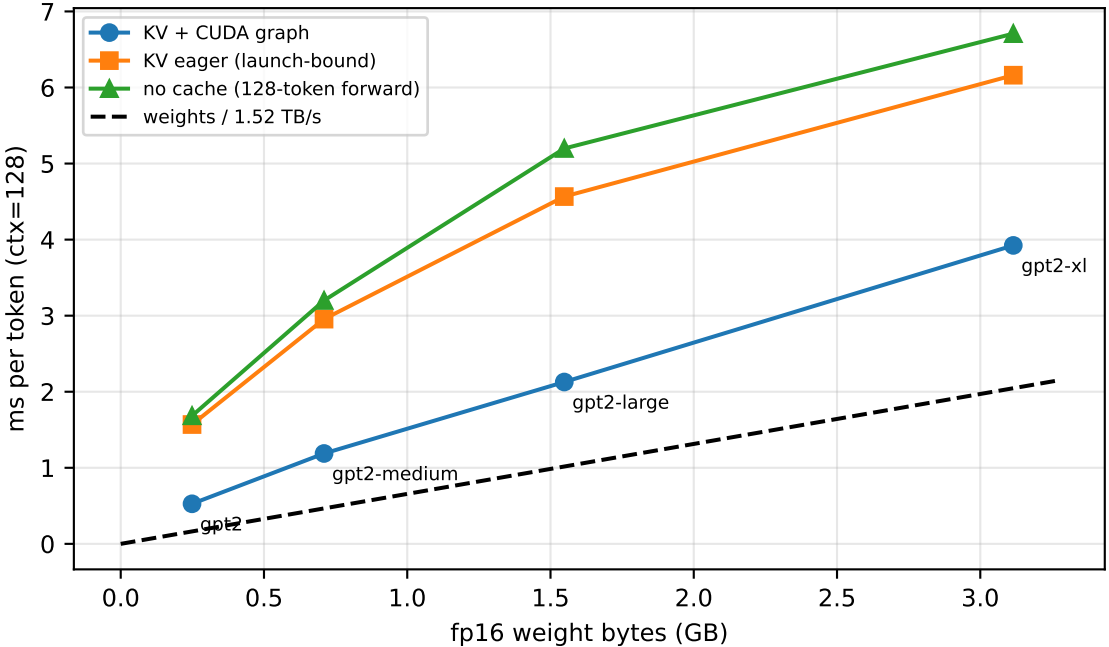
decode latency vs model size
KV + CUDA graph
KV eager (launch-bound)
no cache (128-token forward)
weights / 1.52 TB/s
gpt2
gpt2-medium
gpt2-large
gpt2-xl
ms per token (ctx=128)
fp16 weight bytes (GB)
0
1
2
3
4
5
6
7
0.0
0.5
1.0
1.5
2.0
2.5
3.0