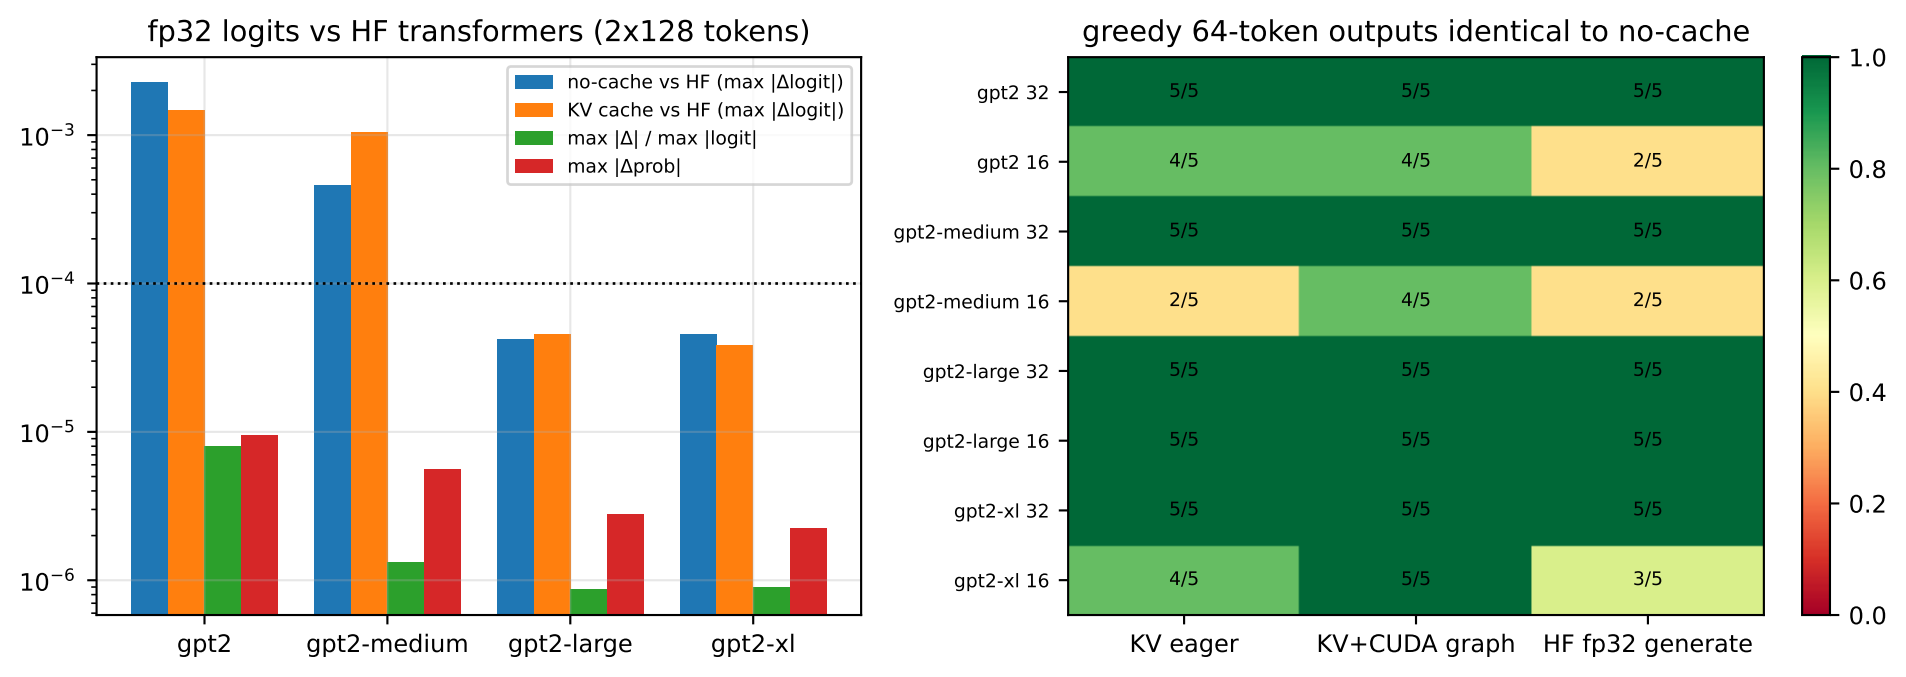
fp32 logits vs HF transformers (2x128 tokens)
no-cache vs HF (max |Δlogit|)
KV cache vs HF (max |Δlogit|)
max |Δ| / max |logit|
max |Δprob|
10^{-3}
10^{-4}
10^{-5}
10^{-6}
gpt2
gpt2-medium
gpt2-large
gpt2-xl
greedy 64-token outputs identical to no-cache
gpt2 32
gpt2 16
gpt2-medium 32
gpt2-medium 16
gpt2-large 32
gpt2-large 16
gpt2-xl 32
gpt2-xl 16
5/5
5/5
5/5
4/5
4/5
2/5
5/5
5/5
5/5
2/5
4/5
2/5
5/5
5/5
5/5
5/5
5/5
5/5
5/5
5/5
5/5
4/5
5/5
3/5
KV eager
KV+CUDA graph
HF fp32 generate
1.0
0.8
0.6
0.4
0.2
0.0